

# Best Practices for Vectorization

Getting ready for Intel® Advanced Vector Extensions 512  
(Intel® AVX-512)

Manel Fernández

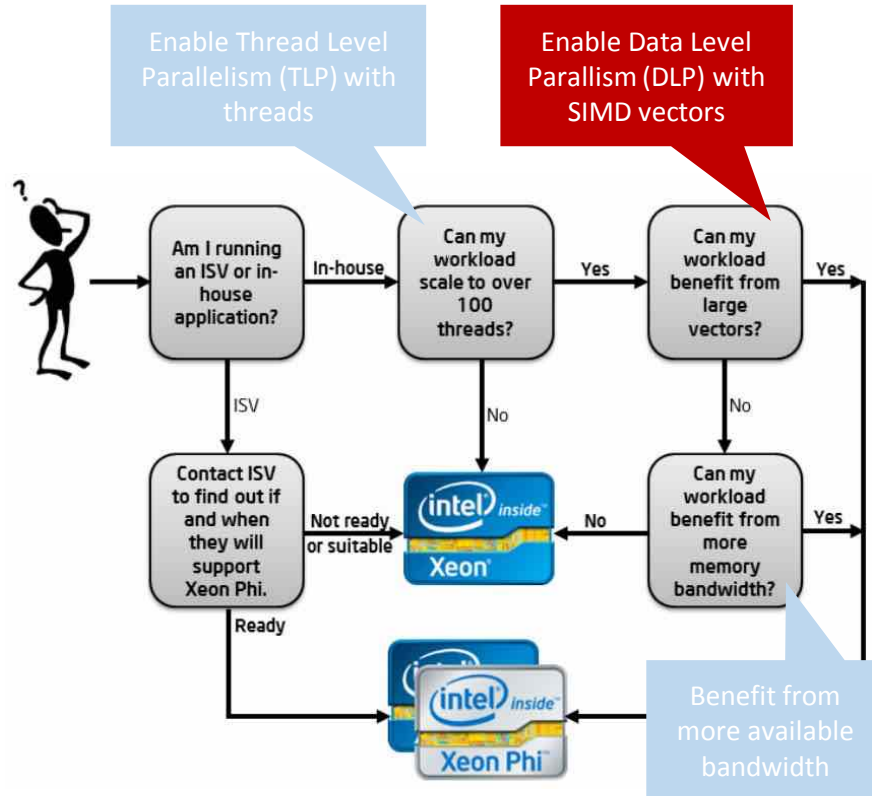
Intel HPC Software Workshop Series 2015

May 6<sup>th</sup> 2015, Munich



# The need for SIMD vectorization

## Is the Intel® Xeon Phi™ coprocessor right for me?



Single thread (ST) performance is limited in today's CPUs

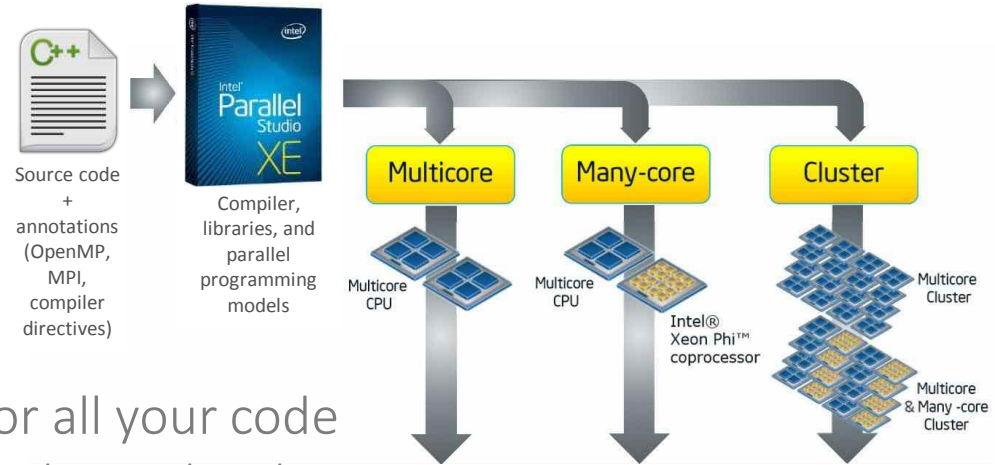
- Clock frequency constraints
- Difficult to discover “near” Instruction level parallelism (ILP) by hardware

More transistors dedicated to exploit “distant” parallelism

- Task level parallelism (TLP)
  - Improves Multi Thread performance (MT)
- Data level parallelism (DLP)
  - Improves Single Thread performance (ST)
  - Enabled by using SIMD vectors

# How to enable SIMD vectorization?

Enabling parallelism with Intel® Parallel Studio XE 2015 tool suite



Single programming model for all your code

- Based on standards: OpenMP/MPI, C/C++/Fortran
- Programmers/tools responsibility to expose DLP/TLP parallelism

Exposing TLP/DLP in your application will benefit today and future Intel® Xeon® processors and Intel® Xeon Phi™ coprocessors

- Including SIMD vectorization on future Intel® AVX-512 products

# Single Instruction Multiple Data (SIMD)

Technique for exploiting DLP on a single thread

- Operate on more than one element at a time
- Might decrease instruction counts significantly

Elements are stored on SIMD registers or *vectors*

Code needs to be *vectorized*


- Vectorization usually on *inner* loops
- Main and *remainder* loops are generated

Scalar loop

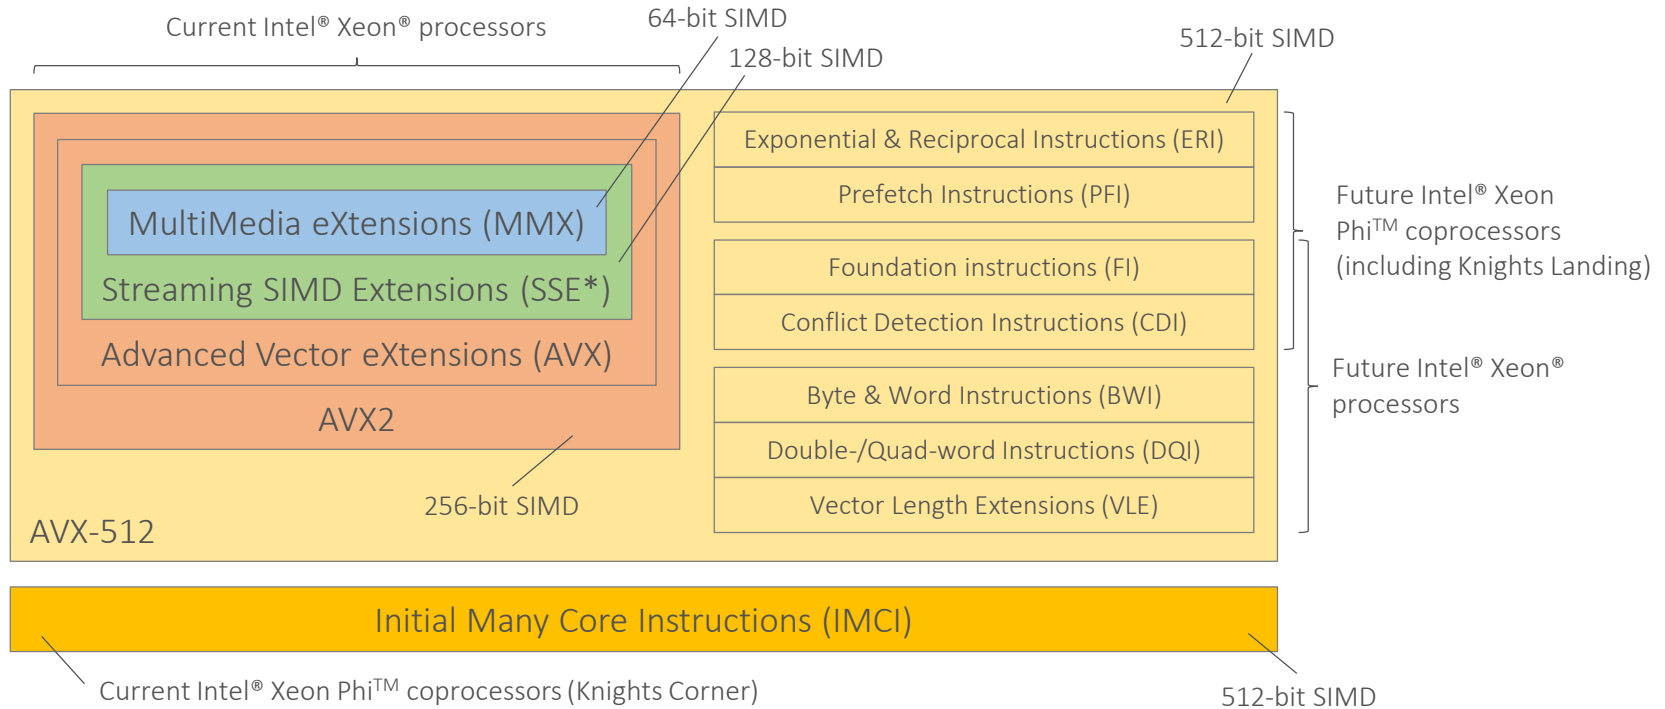
```
for (int i = 0; i < N; i++)  
    c[i] = a[i] + b[i];
```

SIMD loop (4 elements)

```
for (int i = 0; i < N; i += 4)  
    c[i:4] = a[i:4] + b[i:4];
```



# Past, present, and future of Intel SIMD types



For more information about Intel® AVX-512 instructions, check out James Reinders' [initial](#) and [updated](#) post for this topic.

# Intel® AVX2/IMCI/AVX-512 differences

	Intel® Initial Many Core Instructions <b>IMCI</b>	Intel® Advanced Vector Extensions 2 <b>AVX2</b>	Intel® Advanced Vector Extensions 512 <b>AVX-512</b>
Introduction	2012	2013	2015
Products	Knights Corner	Haswell, Broadwell	Knights Landing, future Intel® Xeon® and Xeon® Phi™ products
Register file	SP/DP/int32/int64 data types 32 x 512-bit SIMD registers 8 x 16-bit mask registers	SP/DP/int32/int64 data types 16 x 256-bit SIMD registers No mask registers (instr. blending)	SP/DP/int32/int64 data types 32 x 512-bit SIMD registers 8 x (up to) 64-bit mask
ISA features	Not compatible with AVX*/SSE* No unaligned data support Embedded broadcast/cvt/swizzle MVEX encoding	Fully compatible with AVX/SSE* Unaligned data support (penalty)  VEX encoding	Fully compatible with AVX*/SSE* Unaligned data support (penalty) Embedded broadcast/rounding EVEX encoding
Instruction features	Fused multiply-and-add (FMA) Partial gather/scatter Transcendental support	Fused multiply-and-add (FMA) Full gather	Fused multiply-and-add (FMA) Full gather/scatter Transcendental support (ERI only) Conflict detection instructions PFI/BWI/DQI/VLE (if applies)

Intel® AVX-512 is a major step in unifying the instruction set of Intel® MIC and Intel® Xeon® architecture

# Side effects of SIMD vectorization

## Scalar loop

```
float a[1024], b[1024], c[1024];  
...  
for (int i = 0; i < 1024; i++)  
    c[i] = a[i] + b[i];
```

### Assumptions

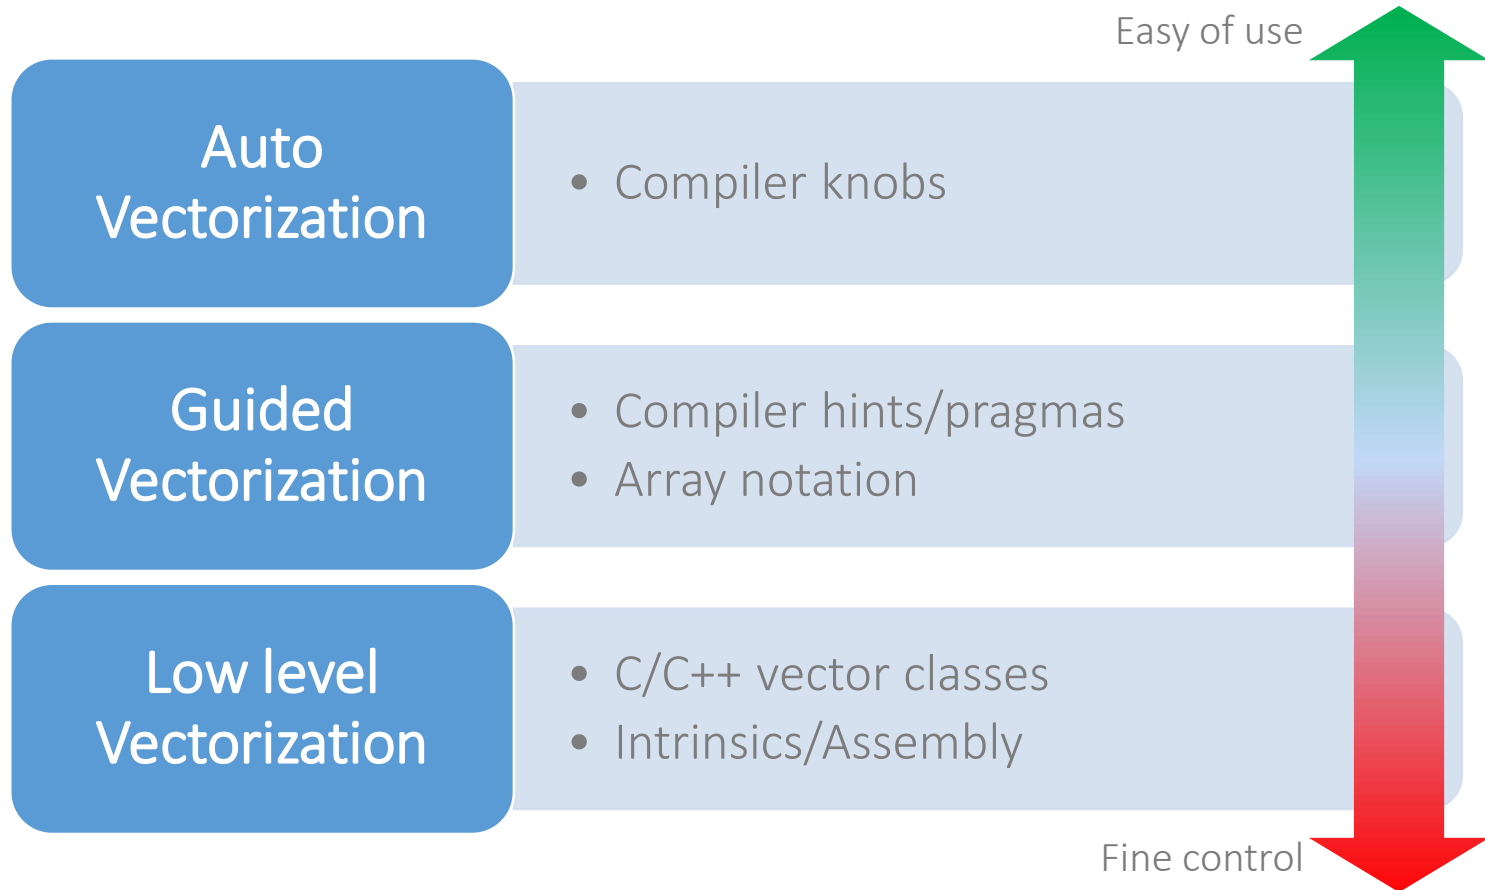
- 64-byte cache lines
- 32-byte (AVX2) and 64-byte (IMCI/AVX-512) SIMD registers
- 4-byte SP elements (float)
- No hardware prefetcher, no ld+op instructions
- Arrays are not cached

#Instructions	Scalar	AVX2 (256-bit)	IMCI AVX-512 (512-bit)
Loads (hit) to a[], b[]	960 + 960	64 + 64	0
Loads (miss) to a[], b[]	64 + 64	64 + 64	64 + 64
SP adds	1024	128	64
Stores to c[]	1024	128	64
<b>Total (Reduction)</b>	<b>4096 (x1)</b>	<b>512 (x8)</b>	<b>256 (x16)</b>

## Observations

- Significant instruction count reduction (up to *vector-length*)
  - IPC decreases, but so does execution time as well
  - Usually translated into speedup
- Compute-bound codes turn into memory-bound codes
  - If code already was memory bound, no benefits at all (other than energy reduction)

# Vectorization on Intel® compilers





# Auto vectorization

Relies on the compiler for vectorization

- No source code changes
- Enabled with `-vec` compiler knob (default in `-O2` and `-O3` modes)

Option	Description
<code>-O0</code>	Disables all optimizations.
<code>-O1</code>	Enables optimizations for speed which are know to not cause code size increase.
<code>-O2/-O</code> (default)	Enables intra-file interprocedural optimizations for speed, including: <ul style="list-style-type: none"><li>• <b>Vectorization</b></li><li>• <b>Loop unrolling</b></li></ul>
<code>-O3</code>	<p>Performs O2 optimizations and enables more aggressive loop transformations such as:</p> <ul style="list-style-type: none"><li>• <b>Loop fusion</b></li><li>• <b>Block unroll-and-jam</b></li><li>• <b>Collapsing IF statements</b></li></ul> <p>This option is recommended for applications that have loops that heavily use floating-point calculations and process large data sets. However, it might incur in slower code, numerical stability issues, and compilation time increase.</p>

Compiler smart enough to apply loop transformations

- It will allow to vectorize more loops

# Vectorization: target architecture options

On which architecture do we want to run our program?

Option	Description
<a href="#"><code>-mmic</code></a>	Builds an application that runs natively on Intel® MIC Architecture.
<a href="#"><code>-xfeature</code></a> <a href="#"><code>-xHost</code></a>	<p>Tells the compiler which processor features it may target, referring to which instruction sets and optimizations it may generate (not available for Intel® Xeon Phi™ architecture). Values for <i>feature</i> are:</p> <ul style="list-style-type: none"> <li>• <b>COMMON-AVX512</b> (includes AVX512 FI and CDI instructions)</li> <li>• <b>MIC-AVX512</b> (includes AVX512 FI, CDI, PFI, and ERI instructions)</li> <li>• <b>CORE-AVX512</b> (includes AVX512 FI, CDI, BWI, DQI, and VLE instructions)</li> <li>• <b>CORE-AVX2</b></li> <li>• <b>CORE-AVX-I</b> (including RDRND instruction)</li> <li>• <b>AVX</b></li> <li>• <b>SSE4.2, SSE4.1</b></li> <li>• <b>ATOM_SSE4.2, ATOM_SSSE3</b> (including MOVBE instruction)</li> <li>• <b>SSSE3, SSE3, SSE2</b></li> </ul> <p>When using <b>-xHost</b>, the compiler will generate instructions for the highest instruction set available on the compilation host processor.</p>
<a href="#"><code>-axfeature</code></a>	Tells the compiler to generate multiple, feature-specific auto-dispatch code paths for Intel® processors if there is a performance benefit. Values for <i>feature</i> are the same described for <code>-xfeature</code> option. Multiple features/paths possible, e.g.: <code>-axSSE2, AVX</code> . It also generates a baseline code path for the default case.

Vectorized code will be different depending on the chosen target architecture

# Auto vectorization: not all loops will vectorize

## Data dependencies between iterations

- Proven Read-after-Write data (i.e., loop carried) dependencies
- Assumed data dependencies
  - Aggressive optimizations (e.g., IPO) might help

RaW dependency

```
for (int i = 0; i < N; i++)  
    a[i] = a[i-1] + b[i];
```

## Vectorization won't be efficient

- Compiler estimates how better the vectorized version will be
- Affected by data alignment, data layout, etc.

Inefficient vectorization

```
for (int i = 0; i < N; i++)  
    a[c[i]] = b[d[i]];
```

## Unsupported loop structure


- While-loop, for-loop with unknown number of iterations
- Complex loops, unsupported data types, etc.
- (Some) function calls within loop bodies
  - Not the case for SVMML functions

Function call within loop body

```
for (int i = 0; i < N; i++)  
    a[i] = foo(b[i]);
```

# Auto vectorization on Intel® compilers

Vectorization breakdown for loop candidates in Polyhedron benchmark suite



[Polyhedron](#) benchmark suite

Intel® Xeon Phi™ 7120A, 61 cores x 4 threads

Intel® Fortran Compiler 15.0.1.14 [-O3 -fp-model fast=2 -align array64byte -ipo -mmic]

# Validating vectorization success

## Generate compiler report about optimizations

- `-qopt-report [=n]`      Generate report (level [1..5], default 2)
- `-qopt-report-file=<fname>`      Optimization report file (stderr, stdout also valid)
- `-qopt-report-phase=<phase>`      Info about opt. phase:

```

LOOP BEGIN at gas_dyn2.f90(193,11) inlined into gas_dyn2.f90(4326,31)
  remark #15300: LOOP WAS VECTORIZED
  remark #15448: unmasked aligned unit stride loads: 1
  remark #15450: unmasked unaligned unit stride loads: 1
  remark #15475: --- begin vector loop cost summary ---
  remark #15476: scalar loop cost: 53
  remark #15477: vector loop cost: 14.870
  remark #15478: estimated potential speedup: 2.520
  remark #15479: lightweight vector operations: 19
  remark #15481: heavy-overhead vector operations: 1
  remark #15488: --- end vector loop cost summary ---
  remark #25456: Number of Array Refs Scalar Replaced In Loop: 1
  remark #25015: Estimate of max trip count of loop=4
LOOP END

```

Vectorized loop

<code>loop</code>	Loop nest optimizations
<code>par</code>	Auto-parallelization
<code>vec</code>	<b>Vectorization</b>
<code>openmp</code>	OpenMP
<code>offload</code>	Offload
<code>ipo</code>	Interprocedural optimizations
<code>pgo</code>	Profile Guided optimizations
<code>cg</code>	Code generation optimizations
<code>tcollect</code>	Trace analyzer (MPI) collection
<code>all</code>	All optimizations (default)

```

LOOP BEGIN at gas_dyn2.f90(2346,15)
  remark #15344: loop was not vectorized: vector dependence prevents vectorization
  remark #15346: vector dependence: assumed OUTPUT dependence between IOLD line 376 and IOLD line 354
  remark #25015: Estimate of max trip count of loop=3000001
LOOP END

```

Non-vectorized loop

# Guided vectorization: disambiguation hints

Get rid of assumed vector dependencies

Assume function arguments won't be aliased

- C/C++: Compile with `-fargument-noalias`

C99 “restrict” keyword for pointers

- Compile with `-restrict` otherwise

```
void v_add(float *restrict c,  
           float *restrict a,  
           float *restrict b)  
{  
    for (int i = 0; i < N; i++)  
        c[i] = a[i] + b[i];  
}
```

Ignore assumed vector dependencies (compiler directive)

- C/C++: `#pragma ivdep`
- Fortran: `!dir$ ivdep`

```
void v_add(float *c, float *a, float *b)  
{  
    #pragma ivdep  
    for (int i = 0; i < N; i++)  
        c[i] = a[i] + b[i];  
}
```

# Some Intel® compiler directives

Directive	Description
<code>distribute, distribute_point</code>	Instructs the compiler to prefer loop distribution at the location indicated.
<code>inline</code>	Instructs the compiler to inline the calls in question.
<code>ivdep</code>	Instructs the compiler to ignore assumed vector dependencies.
<code>loop_count</code>	Indicates the loop count is likely to be an integer.
<code>optimization_level</code>	Enables control of optimization for a specific function.
<code>parallel/noparallel</code>	Facilitates auto-parallelization of an immediately following loop; using keyword <code>always</code> forces the compiler to auto-parallelize; <code>noparallel</code> pragma prevents auto-parallelization.
<code>[no]unroll</code>	Instructs the compiler the number of times to unroll/not to unroll a loop
<code>[no]unroll_and_jam</code>	Prevents or instructs the compiler to partially unroll higher loops and jam the resulting loops back together.
<code>unused</code>	Describes variables that are unused (warnings not generated).
<code>[no]vector</code>	Specifies whether the loop should be vectorised. In case of forcing vectorization that should be according to the given <a href="#">clauses</a> .

# Guided vectorization: `#pragma simd`

Force loop vectorization ignoring **all** dependencies

- Additional [clauses](#) for specify reductions, etc.

SIMD loop

```
void v_add(float *c, float *a, float *b)
{
    #pragma simd
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];
}
```

SIMD function

```
__declspec(vector)
void v_add(float c, float a, float b)
{
    c = a + b;
}

...
for (int i = 0; i < N; i++)
    v_add(C[i], A[i], B[i]);
```

Also supported in OpenMP

- Almost same functionality/syntax
  - Use `#pragma omp simd [clauses]` for SIMD loops
  - Use `#pragma omp declare simd [clauses]` for SIMD functions
- See [OpenMP 4.0 specification](#) for more information



# Intel® compiler directives for vectorization

Directive	Clause	Description
vector	<code>always</code>	Force vectorization even when it might be not efficient.
	<code>[un]aligned</code>	Use [un]aligned data movement instructions for all array vector references.
	<code>[non]temporal (var1[,...])</code>	Do or do not generate non-temporal (streaming) stores for the given array variables. On Intel® MIC architecture, generates a cache-line-evict instruction when the store is known to be aligned.
	<code>[no]vecreminder</code>	Do (not) vectorize the remainder loop when the mail loop is vectorized.
	<code>[no]mask_readwrite</code>	Enables/disables memory speculation causing the generation of [non-]masked loads and stores within conditions.
simd	<code>vectorlength (n1[,...])</code> <code>vectorlengthfor (dtype)</code>	Assume safe vectorization for the given vector length values or data type.
	<code>private (var1[,...])</code> <code>firstprivate (var1[,...])</code> <code>lastprivate (var1[,...])</code>	Which variables are private to each iteration; <i>firstprivate</i> , initial value is broadcasted to all private instances; <i>lastprivate</i> , last value is copied out from the last instance.
	<code>linear (var1:step1[,...])</code>	Letting know the compiler that <i>var1</i> is incremented by <i>step1</i> on every iteration of the original loop.
	<code>reduction (oper: var1[,...])</code>	Which variables are reduction variables with a given operator.
	<code>[no]assert</code>	Warning or error when vectorization fails.
	<code>[no]vecremainder</code>	Do (not) vectorize the remainder loop when the mail loop is vectorized.

# Explicit vectorization with array notation

Express high-level vector parallel array operations

- Valid notation in Fortran since Fortran 90
- Supported in C/C++ by Intel® compiler ([Cilk™ Plus](#)) and GCC 4.9
  - Enabled by default on Intel® compiler, use `-fcilkplus` option on GCC
- No additional modifications to source code
- Most arithmetic and logic operations already overloaded
- Also built-in reducers for array sections

Vectorization becomes explicit

- C/C++ syntax: `array-expression[lower-bound:length[:stride]]`

Samples

```
a[:]      // All elements
a[2:6]    // Elements 2 to 7
a[:,5]    // Column 5
a[0:3:2]  // Elements 0,2,4
```

SIMD function invoked with array notation

```
__declspec(vector)
void v_add(float c, float a, float b)
{
    c = a + b;
}
...
v_add(C[:, A[:, B[:]]];
```

# Improving vectorization: data layout

## Vectorization more efficient with unit strides

- Non-unit strides will generate gather/scatter
- Unit strides also better for data locality
- Compiler might refuse to vectorize

## AoS vs SoA

- Layout your data as Structure of Arrays (SoA)

## Traverse matrices in the right direction

- C/C++: `a[i][:]`, Fortran: `a(:,i)`
- Loop interchange might help
  - Usually the compiler is smart enough to apply it
  - Check compiler optimization report

### Array of Structures vs Structure of Arrays

```
// Array of Structures (AoS)
struct coordinate {
    float x, y, z;
} crd[N];

...
for (int i = 0; i < N; i++)
    ... = ... f(crd[i].x, crd[i].y, crd[i].z);
```

Consecutive elements in memory →

x0 y0 z0 x1 y1 z1 ... x(n-1) y(n-1) z(n-1)

```
// Structure of Arrays (SoA)
struct coordinate {
    float x[N], y[N], z[N];
} crd;

...
for (int i = 0; i < N; i++)
    ... = ... f(crd.x[i], crd.y[i], crd.z[i]);
```

Consecutive elements in memory →

x0 x1 ... x(n-1) y0 y1 ... y(n-1) z0 z1 ... z(n-1)

# Improving vectorization: data alignment

Unaligned accesses might cause significant performance degradation

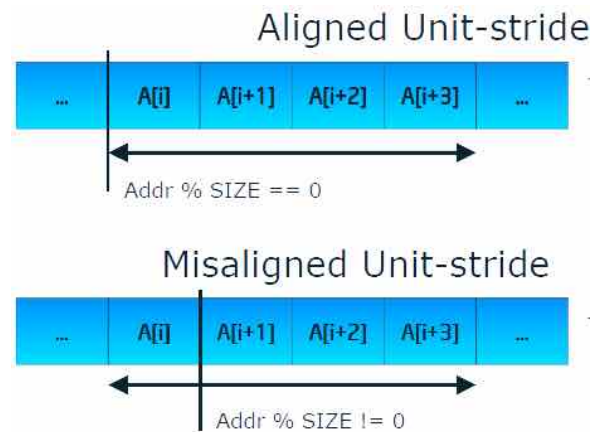
- Two instructions on current Intel® Xeon Phi™ coprocessor
- Might cause “false sharing” problems
  - Consumer/producer thread on the same cache line

Alignment is generally unknown at compile time

- Every vector access is potentially an unaligned access
  - Vector access size = cache line size (64-byte)
- Compiler might “peel” a few loop iterations
  - In general, only one array can be aligned, though

When possible, we have to

- Align our data
- Tell the compiler data is aligned
  - Might not be always the case



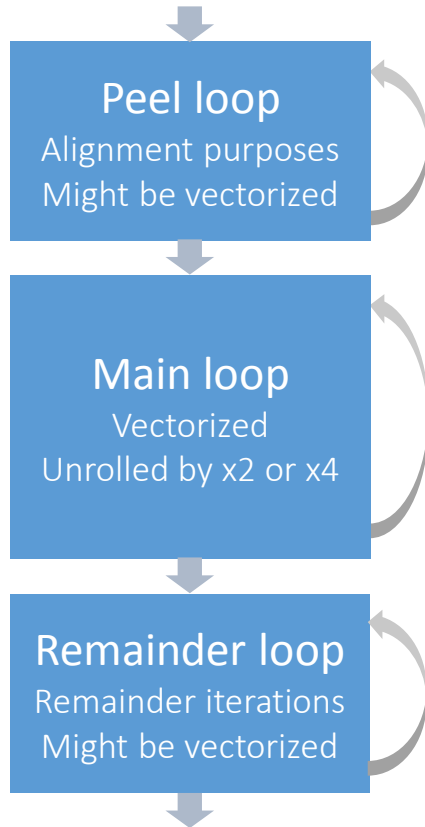
# Improving vectorization: data alignment (cont'd)

How to...	Language	Syntax	Semantics
...align data	C/C++	<code>void* _mm_malloc(int size, int n)</code>	Allocate memory on heap aligned to $n$ byte boundary.
	C/C++	<code>int posix_memalign (void **p, size_t n, size_t size)</code>	
	C/C++	<code>__declspec(align(n)) array</code>	Alignment for variable declarations.
	Fortran (not in common section)	<code>!dir\$ attributes align:n::array</code>	
	Fortran (compiler option)	<code>-alignnbyte</code>	
...tell the compiler about it	C/C++	<code>#pragma vector aligned</code>	Vectorize assuming all array data accessed are aligned (may cause fault otherwise).
	Fortran	<code>!dir\$ vector aligned</code>	
	C/C++	<code>__assume_aligned(array, n)</code>	Compiler may assume array is aligned to $n$ byte boundary.
	Fortran	<code>!dir\$ assume_aligned array:n</code>	

$n=64$  for Intel® Xeon Phi™ coprocessors,  $n=32$  for AVX,  $n=16$  for SSE

Padding might be necessary to guarantee aligned access to matrices

# Vectorization with multi-version loops



```
LOOP BEGIN at gas_dyn2.f90(2330,26)
```

```
<Peeled>
```

```
remark #15389: vectorization support: reference AMAC1U has unaligned access
remark #15381: vectorization support: unaligned access used inside loop body
remark #15301: PEEL LOOP WAS VECTORIZED
```

```
LOOP END
```

```
LOOP BEGIN at gas_dyn2.f90(2330,26)
```

```
remark #25084: Preprocess Loopnests: Moving Out Store
remark #15388: vectorization support: reference AMAC1U has aligned access
remark #15399: vectorization support: unroll factor set to 2
remark #15300: LOOP WAS VECTORIZED
remark #15475: --- begin vector loop cost summary ---
remark #15476: scalar loop cost: 8
remark #15477: vector loop cost: 0.620
remark #15478: estimated potential speedup: 15.890
remark #15479: lightweight vector operations: 5
remark #15488: --- end vector loop cost summary ---
remark #25018: Total number of lines prefetched=4
remark #25019: Number of spatial prefetches=4, dist=8
remark #25021: Number of initial-value prefetches=6
```

```
LOOP END
```

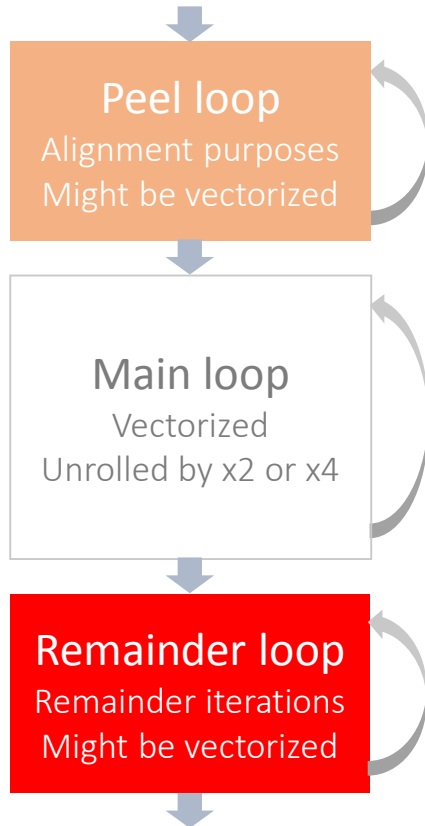
```
LOOP BEGIN at gas_dyn2.f90(2330,26)
```

```
<Remainder>
```

```
remark #15388: vectorization support: reference AMAC1U has aligned access
remark #15388: vectorization support: reference AMAC1U has aligned access
remark #15301: REMAINDER LOOP WAS VECTORIZED
```

```
LOOP END
```

# Improving vectorization: trip count hints



Vectorization can be seen as aggressive unrolling

- Main loop usually unrolled by x2 or x4
- Peel and remainder loop are vectorized with masks
- If trip count is low, vectorization might not be efficient
  - Remainder loop becomes the hotspot

Take a look at remainder loops

- Specify loop trip counts for efficient vectorization
  - `#pragma loop_count (n1, [ n2...])`
  - `#pragma loop_count min(n1), max(n2), avg(n3)`
- Consider padding (Intel® Xeon Phi™ only)
  - Otherwise, remainder loops using gather/scatter loops
  - `-qopt-assume-safe-padding` to avoid it

# Other considerations

Loop tiling/blocking to improve data locality

- Square tiles so elements can be reused

Use streaming loads/stores to save bandwidth

- `#pragma vector [non]temporal(list)`
- `-qopt-streaming-stores=[always|never|auto]`
- `-qopt-streaming-cache-evict [=n]` (Intel® MIC only)

Tune software prefetcher

- `-qopt-prefetch [=n]`
- `-qprefetch-distance=n1[,n2]` (Intel® MIC only)
- `#pragma [no]prefetch [clauses]` (Intel® MIC only)



# Low level (explicit) vectorization

## A.k.a “ninja programming”

Vectorization relies on the programmer with some help from the compiler

SIMD C++ class	Intrinsics	Assembly
<pre>#include &lt;fvec.h&gt;  F32vec4 a,b,c; a = b +c;</pre>	<pre>#include &lt;xmmintrin.h&gt;  __m128 a,b,c; a = _mm_add_ps(b,c);</pre>	<pre>_mm128 a,b,c; __asm {     movaps xmm0,b     movaps xmm1,c     addps xmm0,xmm1     movaps a, xmm0 }</pre>

Might be convenient for low level performance tuning of critical hotspots

Not portable among different SIMD architectures

### Intel Intrinsics Guide

#### Technologies

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☐ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX-512
- ☐ KNC
- ☐ SVML
- ☐ Other

#### Categories

- ☐ Application-Targeted
- ☐ Arithmetic
- ☐ Bit Manipulation
- ☐ Cast
- ☐ Compare
- ☐ Convert
- ☐ Cryptography
- ☒ Elementary Math

The Intel Intrinsics Guide is an interactive reference tool for Intel intrinsic instructions, which are C style functions that provide access to many Intel instructions - including Intel® SSE, AVX, AVX-512, and more - without the need to write assembly code. ✕

sqrt

```
__m512d _mm512_mask_rsqrt14_pd (__m512d src, __mmask8 k, __m512d a) vrsqrt14pd
__m512d _mm512_maskz_rsqrt14_pd (__mmask8 k, __m512d a) vrsqrt14pd
__m512d _mm512_rsqrt14_pd (__m512d a) vrsqrt14pd
```

#### Synopsis

```
__m512d _mm512_rsqrt14_pd (__m512d a)
#include "xmmintrin.h"
Instruction: vrsqrt14pd zmm {k}, zmm
CPUID Flags: AVX512F
```

#### Description

Compute the approximate reciprocal square root of packed double-precision (64-bit) floating-point elements in a, and store the results in dst. The maximum relative error for this approximation is less than 2<sup>-14</sup>.

#### Operation

```
FOR j := 0 to 7
    i := j*64
    dst[i+63:i] := APPROXIMATE(1.0 / SQRT(a[i+63:i]))
ENDFOR
dst[MAX:512] := 0
```

```
__m512 _mm512_mask_rsqrt14_ps (__m512 src, __mmask16 k, __m512 a) vrsqrt14ps
__m512 _mm512_maskz_rsqrt14_ps (__mmask16 k, __m512 a) vrsqrt14ps
__m512 _mm512_rsqrt14_ps (__m512 a) vrsqrt14ps
```

# How to get ready for Intel® AVX-512?

BKM: Start optimizing your application today for current generation of Intel® Xeon® processors and Intel® Xeon™ Phi coprocessors

Tune your AVX-512 kernels on non-existing silicon

- Compile with latest compiler toolchains
  - Intel® compiler (v15.0): `-xCOMMON-AVX512`, `-xMIC-AVX512`, `-xCORE-AVX512`
  - GNU compiler (v4.9): `-mavx512f`, `-mavx512cd`, `-mavx512er`, `-mavx512pf`
- Run Intel® Software Development emulator ([SDE](#))
  - Emulate (future) Intel® Architecture Instruction Set Extensions (e.g. Intel® MPX, ...)
  - Tools available for detailed analysis
    - Instruction type histogram
    - Pointer/misalignment checker
  - Also possible to debug the application while emulated

# Summary

Programmers are mostly responsible of exposing DLP (SIMD) parallelism

Intel® compilers provide sophisticated/flexible support for vectorization

- Auto, guided (assisted), and low-level (explicit) vectorization
- Based on OpenMP standards and specific directives
- Easily portable across different Intel® SIMD architectures

Fine-tuning of generated code is key to achieve the best performance

- Check whether code is actually vectorized
- Data layout, alignment, remainder loops, etc.

Get ready for Intel® AVX-512 by optimizing your application today on current generation of Intel® Xeon® processors and Intel® Xeon™ Phi coprocessors

# Online resources

## Intel® Xeon Phi™

- [Developer portal](#) Programming guides, tools, trainings, case studies, etc.
- [Solutions catalog](#) Existing Intel® Xeon Phi™ solutions for known codes


## Intel® software development tools, performance tuning, etc.

- [Documentation library](#) All available documentation about Intel software
- [Learning lab](#) Learning material with Intel® Parallel Studio XE
- [Performance](#) Resources about performance tuning on Intel hardware
- [Forums](#) Public discussions about Intel SIMD, threading, ISAs, etc.


## Other resources (white papers, benchmarks, case studies, etc.)

- [Go parallel](#) BKM for Intel multi- and many-core architectures
- [Colfax research](#) Publications and material on parallel programming
- [Bayncore labs](#) Research and development activities (WIP)



# Recommended books



*High performance parallelism pearls: multi-core and many-core approaches*, by James Reinders and Jim Jeffers, Morgan Kaufmann, 2014




*Intel® Xeon Phi™ coprocessor high-performance programming*, by Jim Jeffers and James Reinders, Morgan Kaufmann, 2013



*Optimizing HPC applications with Intel® cluster tools*, by Alexander Supalov et al, Apress, 2014

*The software optimization handbook*, by Aart Bik, Intel® press, 2004



*Parallel programming with Intel® Parallel Studio XE*, by Stephen Blair-Chappell and Andrew Stokes, Wrox press, 2012