



CUDA Asynchronous Memory Usage and Execution

Yukai Hung

a0934147@gmail.com


Department of Mathematics

National Taiwan University

Page-Locked Memory


Page-Locked Memory

- Regular pageable and page-locked or pinned host memory
 - use too much page-locked memory reduces system performance




Page-Locked Memory

- Regular pageable and page-locked or pinned host memory
 - copy between page-locked memory and device memory can be performed concurrently with kernel execution for some devices



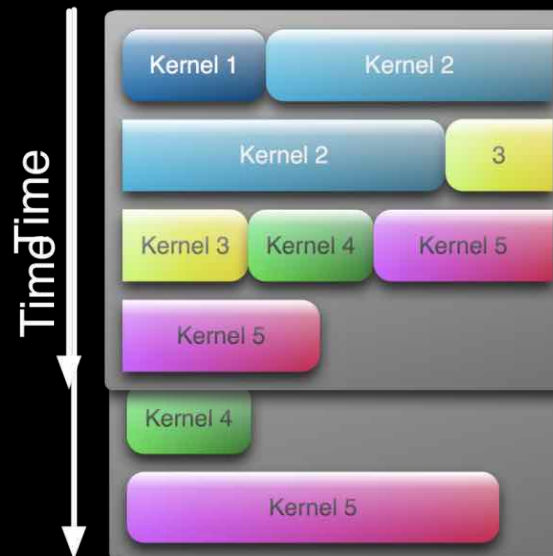
Page-Locked Memory

- Regular pageable and page-locked or pinned host memory
 - copy between page-locked memory and device memory can be performed concurrently with kernel execution for some devices




Page-Locked Memory

- Regular pageable and page-locked or pinned host memory
 - use page-locked host memory can support executing more than one device kernel concurrently for compute capability 2.0 hardware



Page-Locked Memory

- Portable memory
 - the block of page-locked memory is only available for the thread that allocates it by the default setting, use portable memory flag to share the page-locked memory with other threads



Page-Locked Memory

- How to allocate portable memory?

```
float* pointer;
```

```
//allocate host page-locked write-combining memory  
cudaHostAlloc((void**) &pointer, bytes, cudaHostAllocPortable);
```

```
//free allocated memory space  
cudaFreeHost(pointer);
```


Page-Locked Memory

- Write-Combining memory
 - page-locked memory is allocated as cacheable by default
 - page-locked memory can be allocated as write-combining memory by using special flag, which frees up L1 and L2 cache resource usage
- Advantage and disadvantage
 - write-combining memory is not snooped during transfers across bus, which can improve transfer performance by up to 40%
 - reading from write-combining memory from host is slow, which should in general be used for memory that the host only write to

Page-Locked Memory

- How to allocate write-combining memory?

```
float* pointer;
```


```
//allocate host page-locked write-combining memory  
cudaHostAlloc((void**)   
    &pointer, bytes, cudaHostAllocWriteCombined);
```

```
//free allocated memory space  
cudaFreeHost(pointer);
```

1 GB data size	normal	write-combining
host to device	0.533522	0.338092
device to host	0.591750	0.320989


Page-Locked Memory

- Mapped memory
 - the page-locked host memory can be mapped into the address space of the device by passing special flag to allocate memory



Page-Locked Memory

- Mapped memory
 - the page-locked host memory can be mapped into the address space of the device by passing special flag to allocate memory



Page-Locked Memory

- How to allocate mapped memory?

```
float* pointer;
```

```
//allocate host page-locked write-combining memory
```

```
cudaHostAlloc((void**) &pointer, bytes, cudaHostAllocMapped);
```

```
//free allocated memory space
```

```
cudaFreeHost(pointer);
```


Page-Locked Memory

- Check the hardware is support or not?
 - check the hardware properties to ensure it is available for mapping host page-locked memory with device memory

```
cudaDeviceProp deviceprop;  
  
//query the device hardware properties  
//the structure records all device properties  
cudaGetDeviceProperties(&deviceprop,0);  
  
//check the map memory is available or not  
if(!deviceprop.canMapHostMemory)  
printf("cudaError:cannot map host to devicememory\n");
```

Page-Locked Memory

- What is the property structure contents?



Mapped Memory

```
#define size 1048576

int main(int argc, char** argv)
{
    int loop;

    float residual;

    float *h_a, *d_a;
    float *h_b, *d_b;
    float *h_c, *d_c;

    cudaDeviceProp deviceprop;

    //query the device hardware properties
    //the structure records all device properties
    cudaGetDeviceProperties(&deviceprop, 0);

    //check the map memory is available or not
    if(!deviceprop.canMapHostMemory)
        printf("cudaError: cannot map host to device memory\n");
}
```


Mapped Memory

```
//this flag must be set in order to allocate pinned
//host memory that is accessible to the device
cudaSetDeviceFlags(cudaDeviceMapHost);

//allocate host page-locked and accessible to the device memory
//maps the memory allocation on host into cuda device address
cudaHostAlloc((void**) &h_a, sizeof(float) * size, cudaHostAllocMapped);
cudaHostAlloc((void**) &h_b, sizeof(float) * size, cudaHostAllocMapped);
cudaHostAlloc((void**) &h_c, sizeof(float) * size, cudaHostAllocMapped);

//initialize host vectors
for (loop=0; loop<size; loop++)
{
    h_a[loop] = (float) rand() / (RAND_MAX - 1);
    h_b[loop] = (float) rand() / (RAND_MAX - 1);
}

//pass back the device pointer and map with host
cudaHostGetDevicePointer((void**) &d_a, (void*) h_a, 0);
cudaHostGetDevicePointer((void**) &d_b, (void*) h_b, 0);
cudaHostGetDevicePointer((void**) &d_c, (void*) h_c, 0);
```

Mapped Memory

```
//execute device kernel for vector addition
vectorAdd<<<(int)ceil((float)size/256),256s>>>(d_a,d_b,d_c,size);
cudaThreadSynchronize();

//check the result residual value
for(loop=0,residual=0.0;loop<size;loop++)
residual=residual+(h_a[loop]+h_b[loop]-h_c[loop]);

printf("residual value is %f\n",residual);

//free the memory space which must have been returnedd
//by a previous call to cudaMallocHost or cudaHostAlloc
cudaFreeHost(h_a);
cudaFreeHost(h_b);
cudaFreeHost(h_c);

//catch and check cuda error message
if((error=cudaGetLastError())!=cudaSuccess)
printf("cudaError:%s\n",cudaGetErrorString(error));

return 0;
}
```

Mapped Memory

```
__global__ void vectorAdd(float* da,float* db,float* dc,int size)
{
    int index;

    //calculate each thread global index
    index=blockIdx.x*blockDim.x+threadIdx.x;

    if(index<size)
        //each thread computer one component
        dc[index]=da[index]+db[index];

    return;
}
```

Page-Locked Memory

- Several advantages


- there is no need to allocate a block in device memory and copy data between this block and block in host memory, the data transfers are implicitly performed as needed by the kernel
- there is no need to use streams to overlap data transfers with kernel execution, the kernel-originated data transfers overlap with kernel execution automatically
- mapped memory is able to exploit the full duplex of the PCI express bus by reading and writing at the same time, since memory copy only move data in one direction, half duplex

Page-Locked Memory

- Several disadvantages
 - the page-locked memory is shared with host and device, any application must avoid write on the both side simultaneously
 - the atomic functions operating on mapped page-locked memory are not atomic from the point of view of the host or other devices

Page-Locked Memory

- Portable and mapped memory
 - the page-locked host memory can be allocated as both portable and mapped memory, such that each host thread can map the same page-locked memory into different device address



Page-Locked Memory

- Integrated system
 - the mapped page-locked memory is very suitable on integrated system that utilize the a part of host memory as device memory
 - check the integrated field on cuda device properties structure
 - mapped memory is faster, if data only read from or write to global memory once, the coalescing is even more important with mapped memory in order to reduce the data transfer times

Asynchronous Execution

Asynchronous Execution

- Asynchronous execution
 - some functions are supported asynchronous launching in order to facilitate concurrent execution between host and device resource
 - control is returned to the host thread before the work is finished

transfer data between host and device
perform some device kernels
perform some host functions



overlapping

The diagram illustrates the concept of overlapping execution. Three lines of text are listed vertically: 'transfer data between host and device' (yellow), 'perform some device kernels' (yellow), and 'perform some host functions' (green). To the right of these lines, the word 'overlapping' is written in orange. Two white lines originate from the right side of the first two lines and point towards the word 'overlapping', indicating that the first two tasks overlap with the third task.

Asynchronous Execution

- Asynchronous execution

- some functions are supported asynchronous launching in order to facilitate concurrent execution between host and device resource
- control is returned to the host thread before the work is finished

perform device kernel launch

```
kernel<<<blocknum,blocksize,0,stream>>>(...)
```

perform data transfer between host and device

perform data transfer between device and device

```
cudaMemcpyAsync(destination,source,bytes,direction,stream);
```

perform global memory set

Asynchronous Execution

```
#define size 1048576

int main(int argc, char** argv)
{
    int loop;
    int bytes;

    float *h_a, *d_a;
    float *h_b, *d_b;
    float *h_c, *d_c;

    //allocate host page-locked memory
    cudaMallocHost((void**) &h_a, sizeof(float) * size);
    cudaMallocHost((void**) &h_b, sizeof(float) * size);
    cudaMallocHost((void**) &h_c, sizeof(float) * size);

    //allocate device global memory
    cudaMalloc((void**) &d_a, sizeof(float) * size);
    cudaMalloc((void**) &d_b, sizeof(float) * size);
    cudaMalloc((void**) &d_c, sizeof(float) * size);
```

Asynchronous Execution

```
cudaEvent_t stop;  
cudaEvent_t start;  
  
//create an event object which is used to  
//record device execution elapsed time  
cudaCreateEvent(&stop);  
cudaCreateEvent(&start);  
  
//initialize host vectors  
for(loop=0;loop<size;loop)  
{  
    h_a[loop]=(float)rand()/(RAND_MAX-1);;  
    h_b[loop]=(float)rand()/(RAND_MAX-1);;  
}
```

Asynchronous Execution

```
bytes=sizeof(float)*size;

//set time event recorder
cudaEventRecord(start,0);

//copy data from host to device memory asynchronously
cudaMemcpyAsync(d_a,h_a,bytes,cudaMemcpyHostToDevice,0);
cudaMemcpyAsync(d_b,h_b,bytes,cudaMemcpyHostToDevice,0);

//execute device kernel asynchronously
vectorAdd<<<(int)ceil((float)size/256,256,0,0)>>>(d_a,d_,d_c,size);

//copy data from device to host memory asynchronously
cudaMemcpyAsync(h_c,d_c,bytes,cudaMemcpyDeviceToHost,0);

//set time event recorder
cudaEventRecord(stop,0);
```

Asynchronous Execution

```
counter=0;

//increase the counter before the queried
//cuda event has actually been finished
while(cudaEventQuery(stop)==cudaErrorNotReady)
counter=counter+1;

//calculate device execution elapsed time
cudaEventElapsedTime(&elapsed,start,stop);

//check the result residual value
for(loop=0,residual=0.0;loop<size;loop++)
residual=residual+(h_c[loop]-h_a[loop]-h_b[loop]);

printf("counter:%d\n",counter);
printf("residual:%f\n",residual);
```

Asynchronous Execution

```
//free the memory space which must have been returnedd
//by a previous call to cudaMallocHost or cudaHostAlloc
cudaFreeHost(h_a);
cudaFreeHost(h_b);
cudaFreeHost(h_c);

//free the device memory space
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

//free the cuda event object
cudaEventDestroy(stop);
cudaEventDestroy(start);

//catch and check cuda error message
if((error=cudaGetLastError()) != cudaSuccess)
printf("cudaError:%s\n", cudaGetErrorString(error));

return 0;
}
```

Asynchronous Execution

```
__global__ void vectorAdd(float* da,float* db,float* dc,int size)
{
    int index;

    //calculate each thread global index
    index=blockIdx.x*blockDim.x+threadIdx.x;


    if(index<size)
        //each thread computer one component
        dc[index]=da[index]+db[index];

    return;
}
```


Stream

Stream

- Stream
 - applications manage concurrency through stream
 - a stream is a sequence of commands that execute in order
 - all device requests made from the host code are put into a queues




Stream

- How to create a stream?

```
cudaStream_t stream;  
  
//create an asynchronous new stream  
cudaStreamCreate(&stream);  
  
//destroy stream  
cudaStreamDestroy(stream);
```

Stream

- Stream
 - different streams may execute their commands or host requests out of order with respect to one another or concurrently, but the same stream is still a sequence of commands that execute in order



Stream

```
#define snum 10
#define size 1048576

int main(int argc, char** argv)
{
    int loop;
    int bytes;

    float *h_a, *d_a;
    float *h_b, *d_b;
    float *h_c, *d_c;

    cudaStream_t stream[snum];

    //create new asynchronous stream
    //which acts as device work queue
    for (loop=0; loop<snum; loop++)
        cudaStreamCreate(&stream[loop]);
```

Stream

```
//allocate host page-locked memory
cudaMallocHost((void**) &h_a, sizeof(float) * size * snum);
cudaMallocHost((void**) &h_b, sizeof(float) * size * snum);
cudaMallocHost((void**) &h_c, sizeof(float) * size * snum);

//allocate device global memory
cudaMalloc((void**) &d_a, sizeof(float) * size * snum);
cudaMalloc((void**) &d_b, sizeof(float) * size * snum);
cudaMalloc((void**) &d_c, sizeof(float) * size * snum);

//initialize host vectors
for(loop=0; loop<size*snum; loop++)
{
    h_a[loop] = (float) rand() / (RAND_MAX - 1);
    h_b[loop] = (float) rand() / (RAND_MAX - 1);
}
```

Stream

```
//put all the works into default stream
//executes all works by using noly one stream
for(loop=0;loop<snum;loop++)
{
    bytes=sizeof(float)*size;

    sp1=h_a+loop*size; dp1=d_a+loop*size;
    sp2=h_b+loop*size; dp2=d_b+loop*size;
    sp3=d_c+loop*size; dp3=h_c+loop*size;

    //copy data from host to device memory asynchronously
    cudaMemcpyAsync(dp1,sp1,bytes,cudaMemcpyHostToDevice,0);
    cudaMemcpyAsync(dp2,sp2,bytes,cudaMemcpyHostToDevice,0);

    //execute device kernel asynchronously
    kernel<<<blocknum,blocksize,0,0>>>(d_a,d_b,d_c,size);

    //copy data from device to host memory asynchronously
    cudaMemcpyAsync(dp3,sp3,bytes,cudaMemcpyDeviceToHost,0);
}

//wait until the stream is finished
cudaThreadSynchronize();
```

Stream

```
//put all the works into different asynchronous streams
//each stream only executes three copies and one kernel
for(loop=0;loop<snum;loop++)
{
    bytes=sizeof(float)*size;

    sp1=h_a+loop*size; dp1=d_a+loop*size;
    sp2=h_b+loop*size; dp2=d_b+loop*size;
    sp3=d_c+loop*size; dp3=h_c+loop*size;

    //copy data from host to device memory asynchronously
    cudaMemcpyAsync(dp1,sp1,bytes,cudaMemcpyHostToDevice,stream[loop]);
    cudaMemcpyAsync(dp2,sp2,bytes,cudaMemcpyHostToDevice,stream[loop]);

    //execute device kernel asynchronously
    kernel<<<blocknum,blocksize,0,stream[loop]>>>(d_a,d_b,d_c,size);

    //copy data from device to host memory asynchronously
    cudaMemcpyAsync(dp3,sp3,bytes,cudaMemcpyDeviceToHost,stream[loop]);
}

//wait until all stream are finished
cudaThreadSynchronize();
```


Stream

```
//free the memory space which must have been returnedd
//by a previous call to cudaMallocHost or cudaHostAlloc
cudaFreeHost(h_a);
cudaFreeHost(h_b);
cudaFreeHost(h_c);

//free the device memory space
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

//free the asynchronous streams
for(loop=0;loop<snum;loop++)
    cudaStreamDestroy(stream[loop]);

return 0;
}
```

Stream

```
__global__ void vectorAdd(float* da,float* db,float* dc,int size)
{
    int loop;
    int index;

    volatile float temp1;
    volatile float temp2;

    //calculate each thread global index
    index=blockIdx.x*blockDim.x+threadIdx.x;

    if(index<size)
    for(loop=0;loop<iteration;loop++)
    {
        temp1=da[index];
        temp2=db[index];
        dc[index]=temp1+temp2;
    }

    return;
}
```

Stream

- How about the performance ?

	Fermi C2050	Tesla C1060
single stream	64.096382	180.179825
multiple stream	31.996338	166.010757

Stream

- Stream controlling

`cudaThreadSynchronize()`

called in the end to make sure all streams are finished before proceeding further, it forces the runtime to wait until all device tasks or commands in all asynchronous streams have completed

`cudaStreamSynchronize()`

force the runtime to wait until all preceding device tasks or host commands in one specific stream have completed

`cudaStreamQuery()`

provide applications with a way to know if all preceding device tasks or host commands in a stream have completed

Stream


- Stream controlling

`cudaStreamDestroy()`

wait for all preceding tasks in the give stream to complete before destroying the stream and returning control to the host thread, which is blocked until the stream finished all commands or tasks

Stream

- Overlap of data transfer and kernel execution
transfer data between host page-locked memory and device memory and kernel execution can be performed concurrently




Stream

- Overlap of data transfer and kernel execution
transfer data between host page-locked memory and device memory and kernel execution can be performed concurrently
any application may query the hardware capability by calling the device manage function and checking the property flag

```
cudaDeviceProp deviceprop;  
  
//query the device hardware properties  
//the structure records all device properties  
cudaGetDeviceProperties(&deviceprop,0);  
  
//check the overlapping is available or not  
if(!deviceprop.deviceOverlap)  
printf("cudaError:cannot overlap kernel and transfer\n");
```

Stream

- Concurrent kernel execution
some hardware can execute multiple kernels concurrently



Stream

- Concurrent kernel execution

some hardware can execute multiple kernels concurrently

any application may query the hardware capability by calling the device manage function and checking the property flag

```
cudaDeviceProp deviceprop;  
  
//query the device hardware properties  
//the structure records all device properties  
cudaGetDeviceProperties(&deviceprop, 0);  
  
//check the concurrent kernels is available or not  
if(!deviceprop.concurrentKernels)  
printf("cudaError:cannot use concurrent kernels\n");
```


Stream

- Concurrent kernel execution
 - some hardware can execute multiple kernels concurrently
 - any application may query the hardware capability by calling the device manage function and checking the property flag
 - the kernels that may use many textures or registers or shared memory are less likely to execute with other kernels concurrently

Stream


- Concurrent data transfer

some devices can perform copy data from page-locked memory to device memory with copy data from device memory to host page-locked memory concurrently



Stream

- Streams on the matrix-matrix multiplication



Event

Event

- Event

the runtime provides a ways to closely monitor the device progress by letting the program record events at any point in the program

```
cudaEvent_t event1;  
cudaEvent_t event2;  
  
//create and initialize event  
cudaEventCreate(&event1);  
cudaEventCreate(&event2);  
  
//insert event recorder into stream  
cudaEventRecord(event1, stream);  
cudaEventRecord(event2, stream);  
  
//destroy created event recorder  
cudaEventDestroy(event1);  
cudaEventDestroy(event2);
```

Event

- Event

`cudaEventSynchronize()`

this function blocks until the event has actually been recorded ,
since the event recorder is an asynchronous method


`cudaEventQuery()`

provide any applications with a way to know if one specific event
recorder in the stream have completed , which returns `cudaSuccess`

Multi-GPU


Multi-GPU

- GPU can not share global memory
 - one GPU can not access another GPU's memory directly
 - application code is responsible for moving data between GPUs



Multi-GPU

- A host thread can maintain one context at a time
 - need as many host threads as GPUs to maintain all device
 - multiple host threads can establish context with the same GPU hardware driver handles time-sharing and resource partitioning



Multi-GPU

- Device management calls

`cudaGetDeviceCount()`

returns the number of devices on the current system with compute capability greater or equal to 1.0, that are available for execution

`cudaSetDevice()`

set the specific device on which the active host thread executes the device code. If the host thread has already initialized the cuda runtime by calling non-device management runtime functions, returns error must be called prior to context creation, fails if the context has already been established, one can force the context creation with `cudaFree(0)`

`cudaGetDevice()`

returns the device on which the active host thread executes the code

Multi-GPU

- Device management calls

`cudaThreadExit()`

explicitly clean up all runtime-related resource associated with the calling host thread, any subsequent calls reinitializes the runtime, this function is implicitly called on the host thread exit

`cudaGetDeviceProperties()`

returns the properties of one specific device, this is useful when a system contains different devices in order to choose best devices

● Reference

- Mark Harris <http://www.markmark.net/>
- Wei-Chao Chen <http://www.cs.unc.edu/~ciao/>
- Wen-Mei Hwu <http://impact.crhc.illinois.edu/people/current/hwu.php>