

Experiences in the Land of Virtual Abstractions

Galen Hunt

Principal Researcher

Microsoft Research Operating Systems Group



Why do we all love (hardware) VMs?

- Three reasons:
 - **Compatibility**
 - I can install my application in a VM with the OS it requires and never have to worry about it breaking again.
 - **Security**
 - I can download even the most malignant code from the internet, run it in a VM, and my system's integrity isn't lost.
 - **Continuity**
 - I can start an application (in a VM) on one computer and move it to another, or reboot the computer, and it still runs.

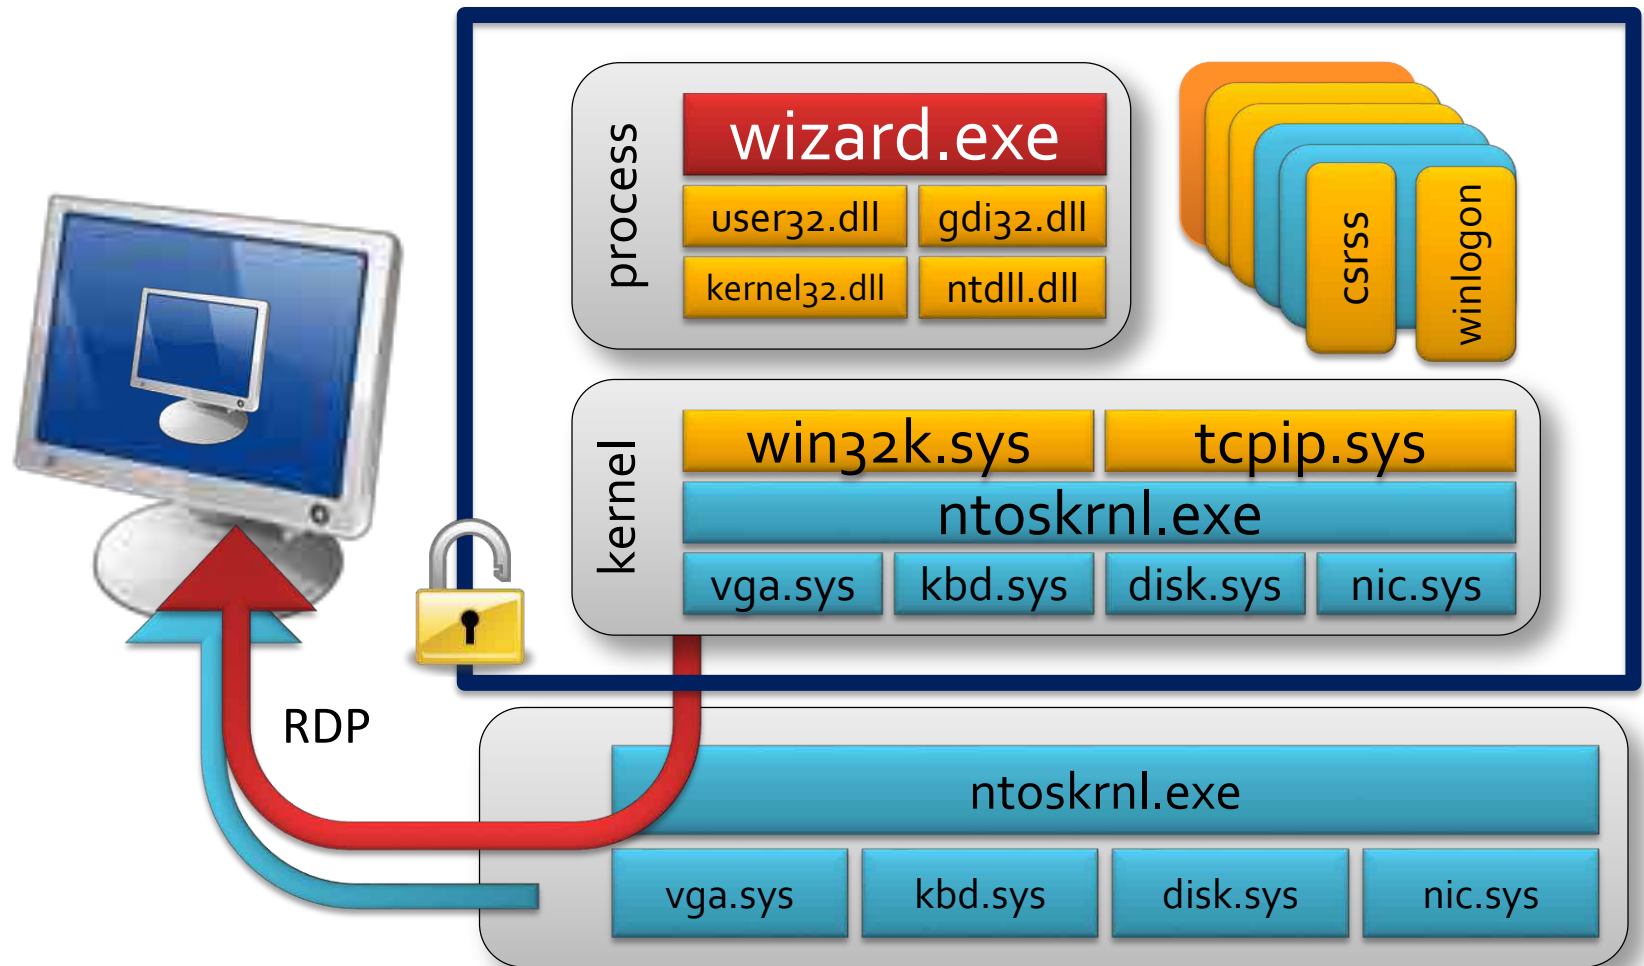
Fantastic Disruption

- VMMs changed server computing forever:
 - server consolidation
 - cloud computing
- Why not desktop and mobile computing?
 - VMs have huge memory and disk overheads
 - “XP mode in Win7” : 1GB VHD, 256MB RAM
 - “Win7 VM” : 4GB+ VHD, \geq 512MB RAM

The Father's Dilemma



Do we need to duplicate the full OS in every VM?



What might be the alternative abstractions?

- Windows abstractions at top of Win32 API?
 - too large to be practical for self-contained apps: 250K+ APIs, registry, etc.
 - not (completely) *compatible* across OS versions
 - not *secure* against untrusted applications
 - not *continuous* as users move from one device to another

What properties do the abstractions need?

- (or what is it that is so “magical” about the VM hardware ABI?)
 - Stable
 - Formalized
 - Closed
 - Transparent (stateless)
- So, ... can we make a higher-level ABI that has these properties?

Drawbridge Abstractions & ABI

Memory Management Primitives (3)

- DkVirtualMemoryAlloc
- DkVirtualMemoryFree
- DkVirtualMemoryProtect

Threading Primitives (16)

- DkThreadCreate¹
- DkThreadDelayExecution
- DkThreadYieldExecution
- DkThreadExit
- DkThreadGetParameter
- DkThreadRaiseException
- DkNotificationEventCreate
- DkSynchronizationEventCreate
- DkSemaphoreCreate
- DkSemaphoreRelease
- DkSemaphorePeek
- DkEventSet
- DkEventClear
- DkEventPeek
- DkObjectsWaitAny
- DkAbortEventRegister¹

Child Process Primitives (3)

- DkProcessCreate¹
- DkProcessGetExitCode¹
- DkProcessExit¹

I/O Stream Primitives (14)

- DkStreamOpen¹
- DkStreamRead²
- DkStreamWrite²
- DkStreamMap²
- DkStreamMapPeBinary²
- DkStreamUnmap²
- DkStreamSetLength²
- DkStreamFlush²
- DkStreamDelete¹
- DkStreamGetEvent²
- DkStreamRename¹
- DkStreamEnumerateChildren¹
- DkStreamAttributesQuery¹
- DkStreamAttributesQueryByHandle²

Other Primitives (9)

- DkSystemTimeQuery
- DkRandomBitsRead
- DkInstructionCacheFlush
- DkObjectReference
- DkObjectClose²
- DkInputEventRead¹
- DkFrameBufferExport¹
- DkFrameBufferNotifyUpdate¹
- DkDebugStringPrint¹

Upcalls (3)

- LibOsInitialize
- LibOsThreadStart
- LibOsExceptionDispatch

Files/Storage (1)

file:

Console Redirection (4)

null:
stderr:
stdin:
stdout:

Named Pipes (2)

pipe.client:
pipe.server:

TCP/IP Stack (4)

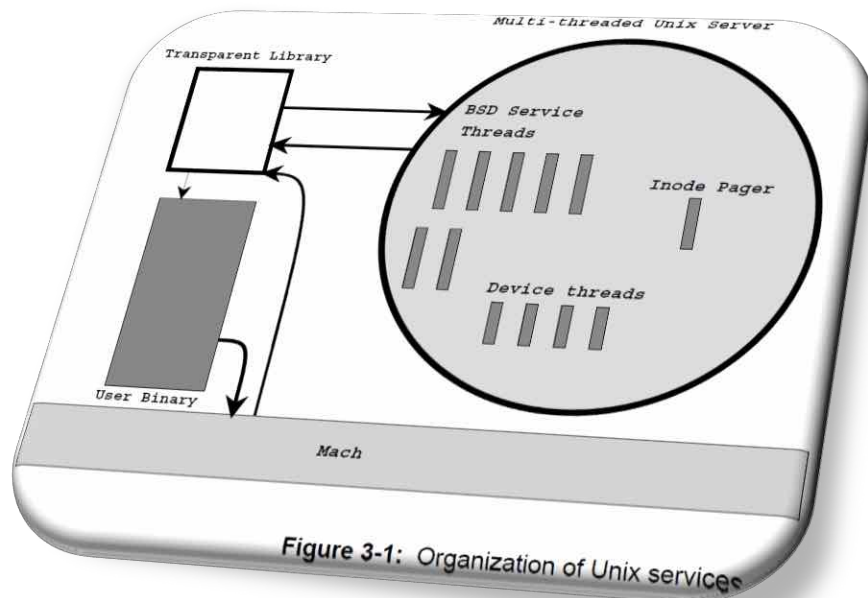
dns:
tcp.client:
tcp.server:
tcp:

HTTP.SYS (2)

http.application:
http.server:

Learning from the past...

“Since March of 1989 we have had running at CMU a computing environment in which the functions of a traditional Unix system are cleanly divided into two parts: facilities which manage the hardware resources of a computer system (such as CPU, I/O and memory) and support for higher-level resource abstractions used in the building of application ...”

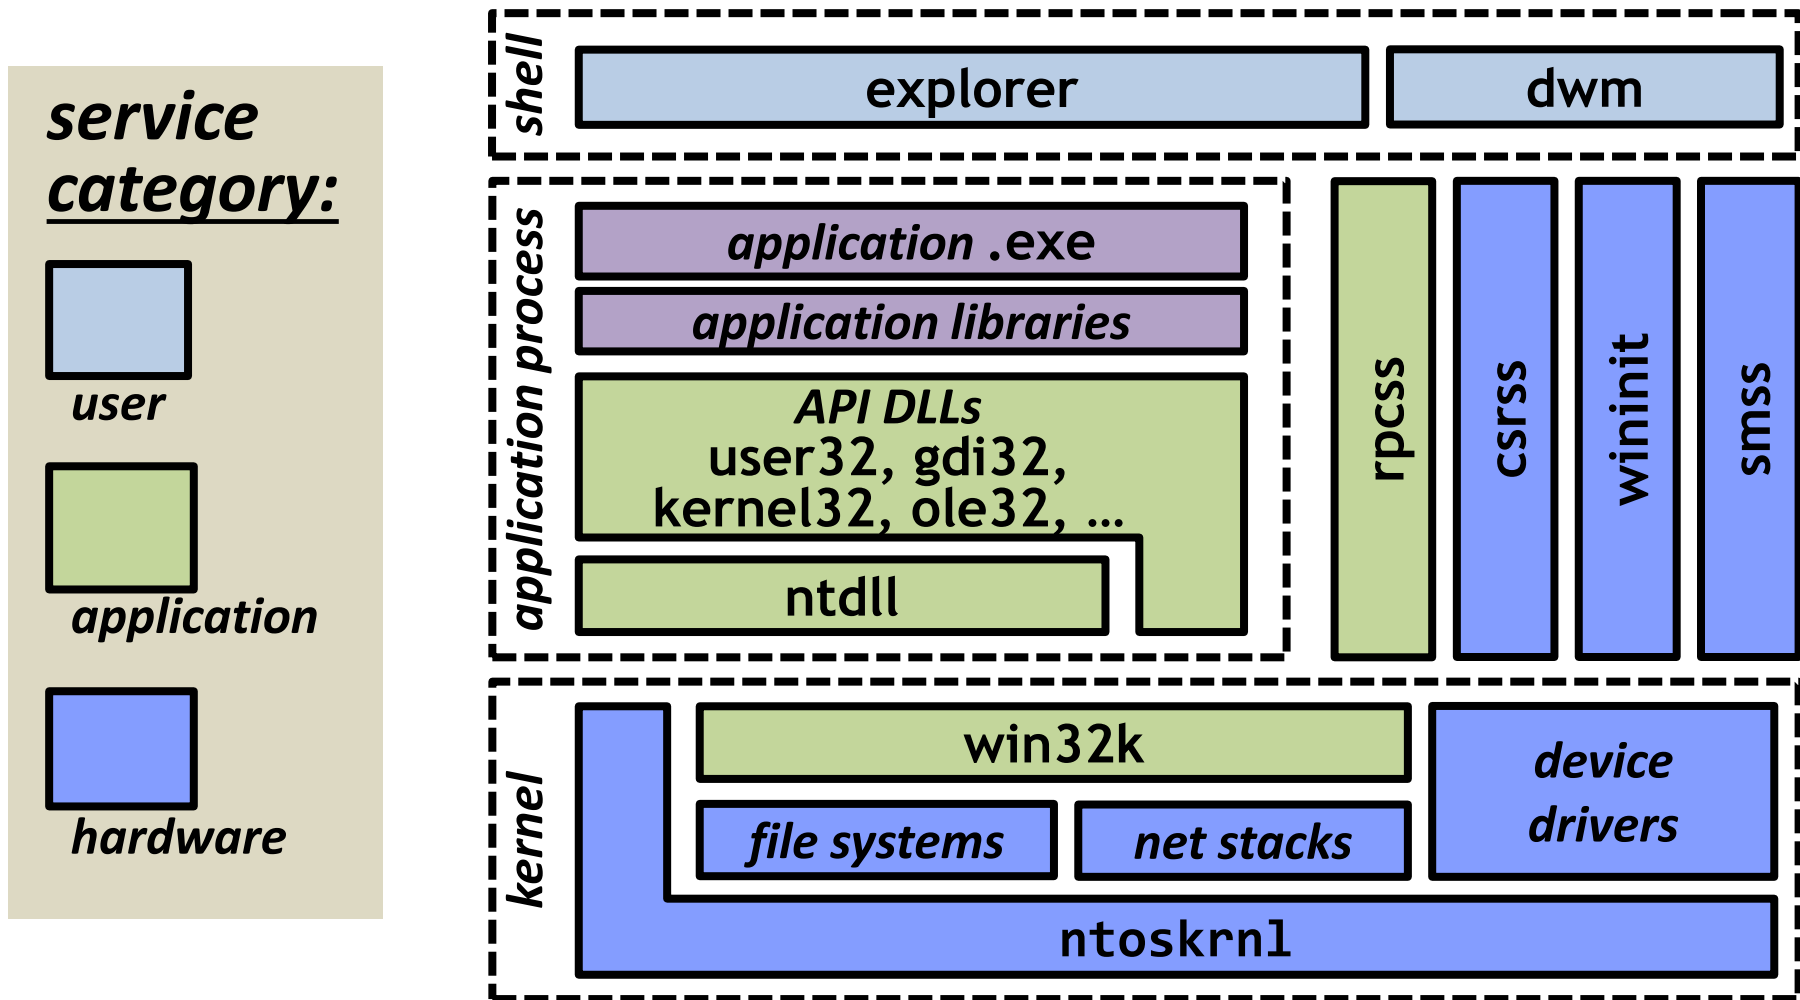


“UNIX as an Application Program”

David B. Golub, Randall W. Dean,
Alessandro Forin, and Richard F. Rashid

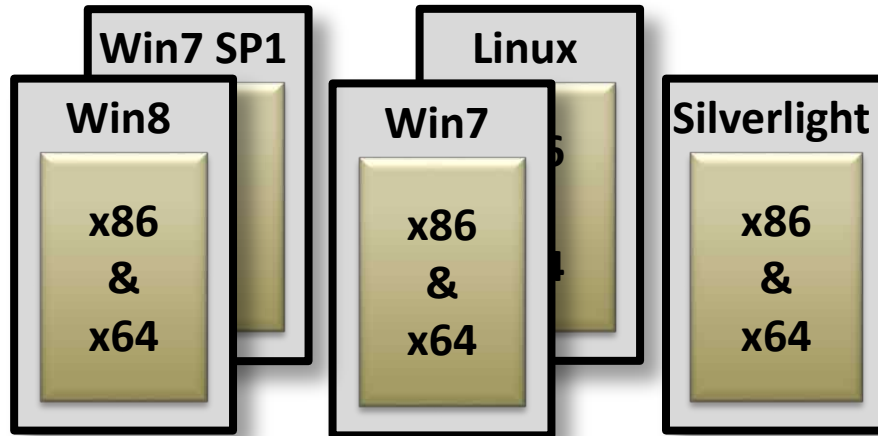
Proc. of the 1990 Summer USENIX Technical Conference

Three categories of services in an OS



Compatibility

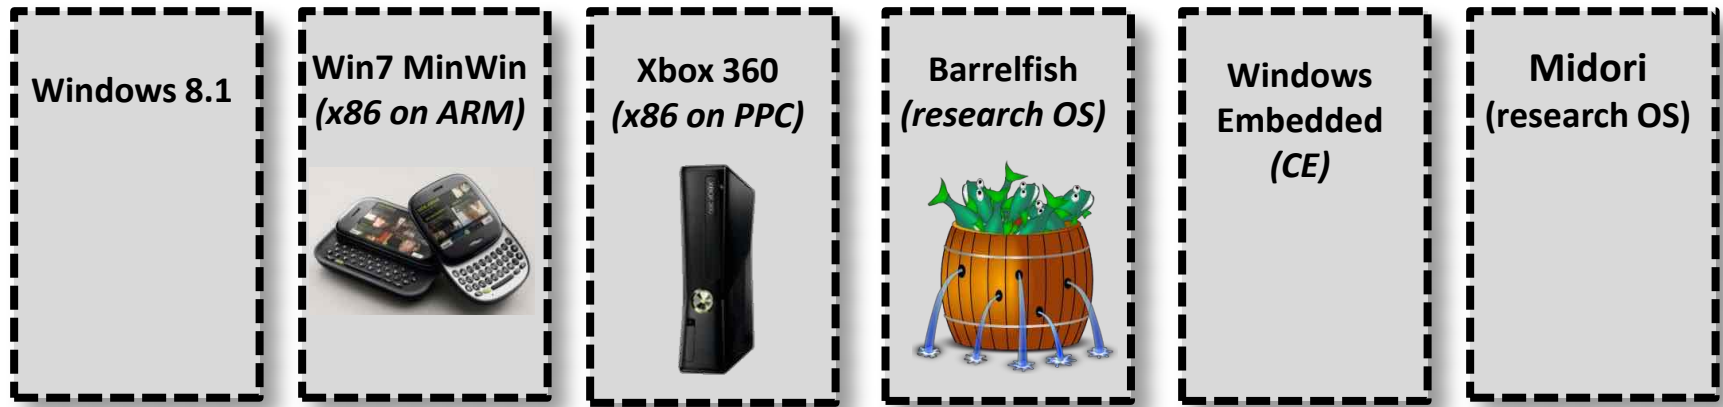
- Existing **library** OS implementations:

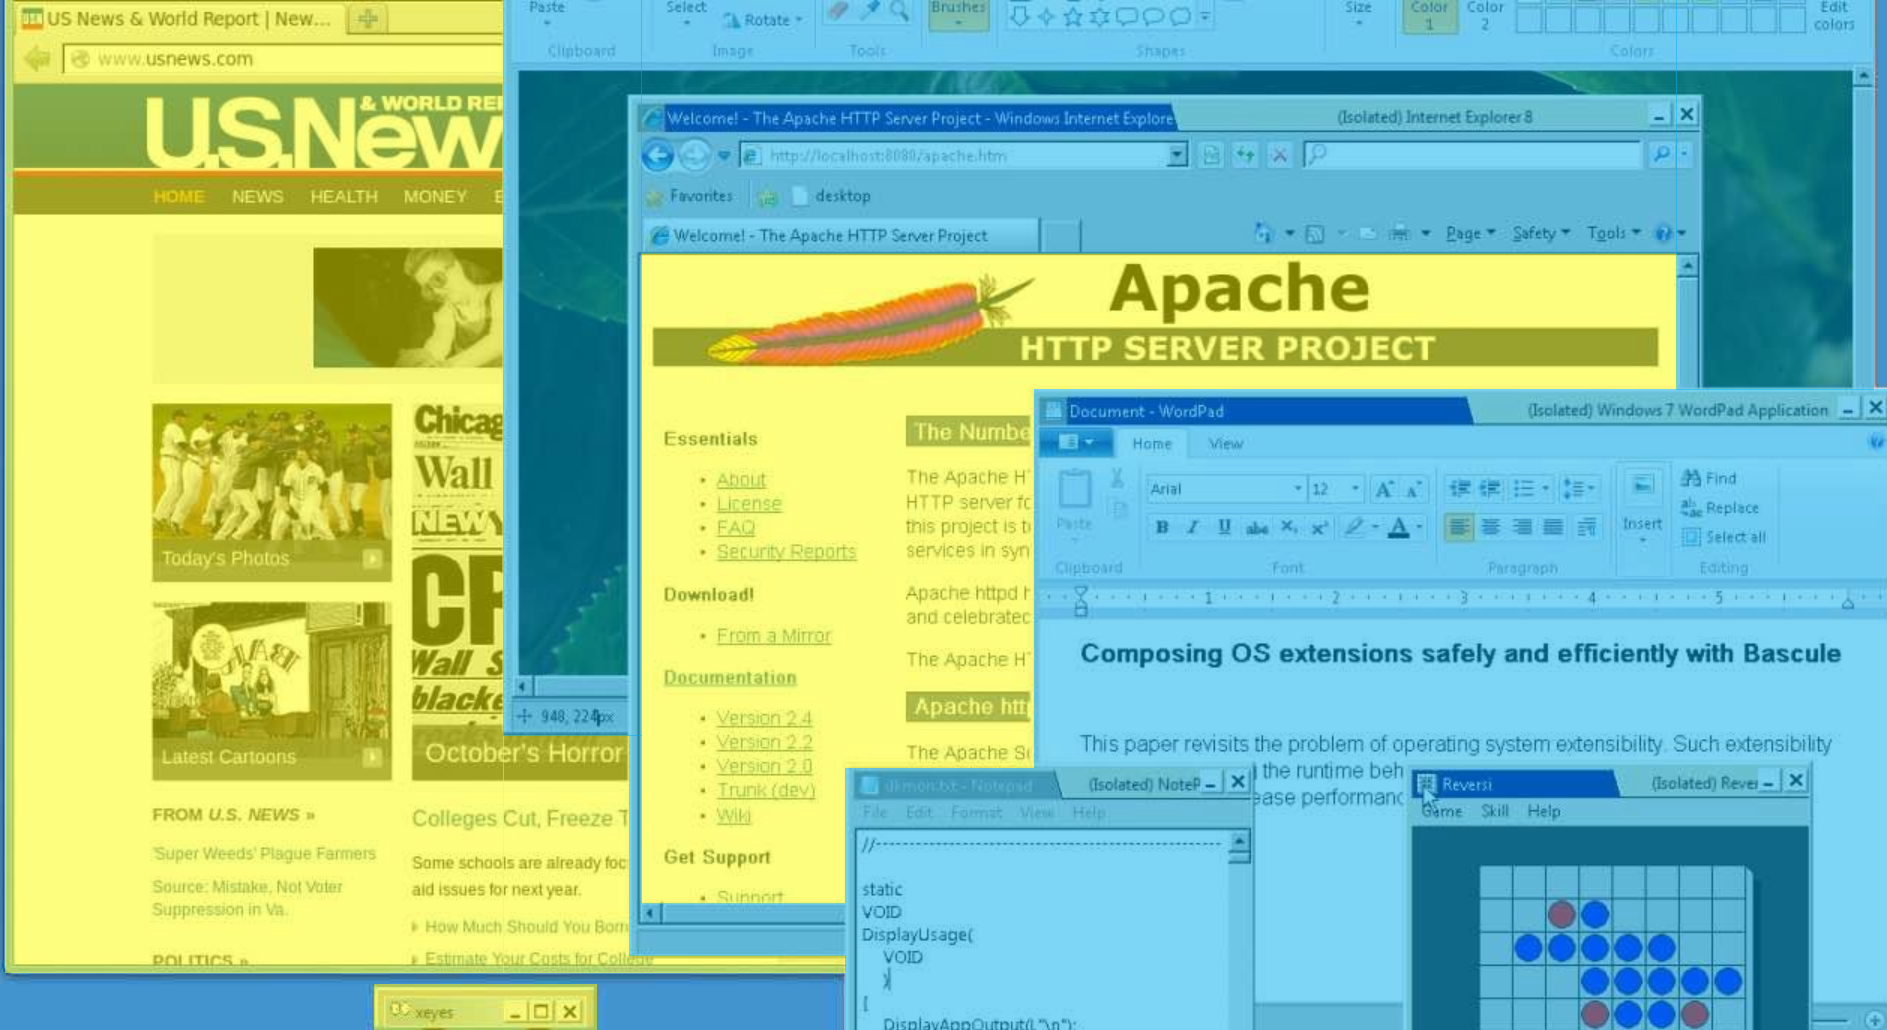


- Full **host** implementations:



- Prototype **host** implementations (*past, present, and in-progress*):





Windows and Linux LibOSes

Linux apps: Firefox, Apache, xeyes

Windows apps: Paint, IE, WordPad, Notepad, Reversi

Security: What's the Thread Model?

- Traditional OS: “**Enterprise Multitenancy**”
 - Invite your employees
 - after full authentication
 - to run the applications you choose (or that your OS vendor vets)
 - on some subset of your computers
- VMs & Drawbridge: “**Hostile Multitenancy**”
 - Invite “organized crime”
 - with complete anonymity (in name & number)
 - to run any code they choose
 - on the same servers as your most valuable customers

Security

Administrator: cmd

Drawbridge

```
[c:\1191]
>monitor\x64\dkmon.exe fontexp.exe corrupt.ttf
Microsoft Drawbridge Console Host [Version 1.0.1191.0]
[Monitor] Writable folder: [C:\Users\Administrator\AppData\Roaming\Microsoft\Drawbridge\Sandboxes\89f54332-d6ef-4ed2-9b8d-ca80b3399516\Content]
[Monitor] Scratch folder: [C:\Users\Administrator\AppData\Roaming\Microsoft\Drawbridge\Sandboxes\89f54332-d6ef-4ed2-9b8d-ca80b3399516\Scratch]
[Monitor] Driver version: [1.0.1191.0] (retail build)
[Monitor] Loaded package: [NT.6.2.9200.1191.0.Picoprocess (x64)]
[Monitor] Loaded package: [WindowsGUI.6.2.9200.0.0.Stub (x64)]
[Monitor] Loaded package: [WindowsGUI.6.2.9200.0.0.Stub (x64)]
[Monitor] Loaded package: [fontexp.exe.0.0.0.0]
[Monitor] Loaded package: [Windows.6.2.9200.0.0.0]
[Monitor] Loaded package: [Drawbridge.1.0.1191.0]
[Monitor] Loaded package: [WindowsCertification]
[Monitor] Loaded package: [InternetExplorer]
[Monitor] Launching app: [fontexp.exe]
[Monitor:FATAL] Exception: 0xc0000005 (!! Last error: 0, 0x0, 0x0)
[Monitor:ERROR] A fatal error occurred in the application. The application will be terminated. Code: 0xc0000005.
[Monitor] Terminating process...
Loading from corrupt.ttf
[Monitor:WARNING] Application 'fontexp.exe' code integrity failed with error 0xc0000005.
[Monitor] Done.

[c:\1191]
>
```

Your PC ran into a problem and needs to restart. We're just collecting some error info, and then we'll restart for you. (0% complete)

If you'd like to know more, you can search online later for this error: SYSTEM_SERVICE_EXCEPTION (win32k.sys)

Native

Continuity

- **Checkpoint Extension**

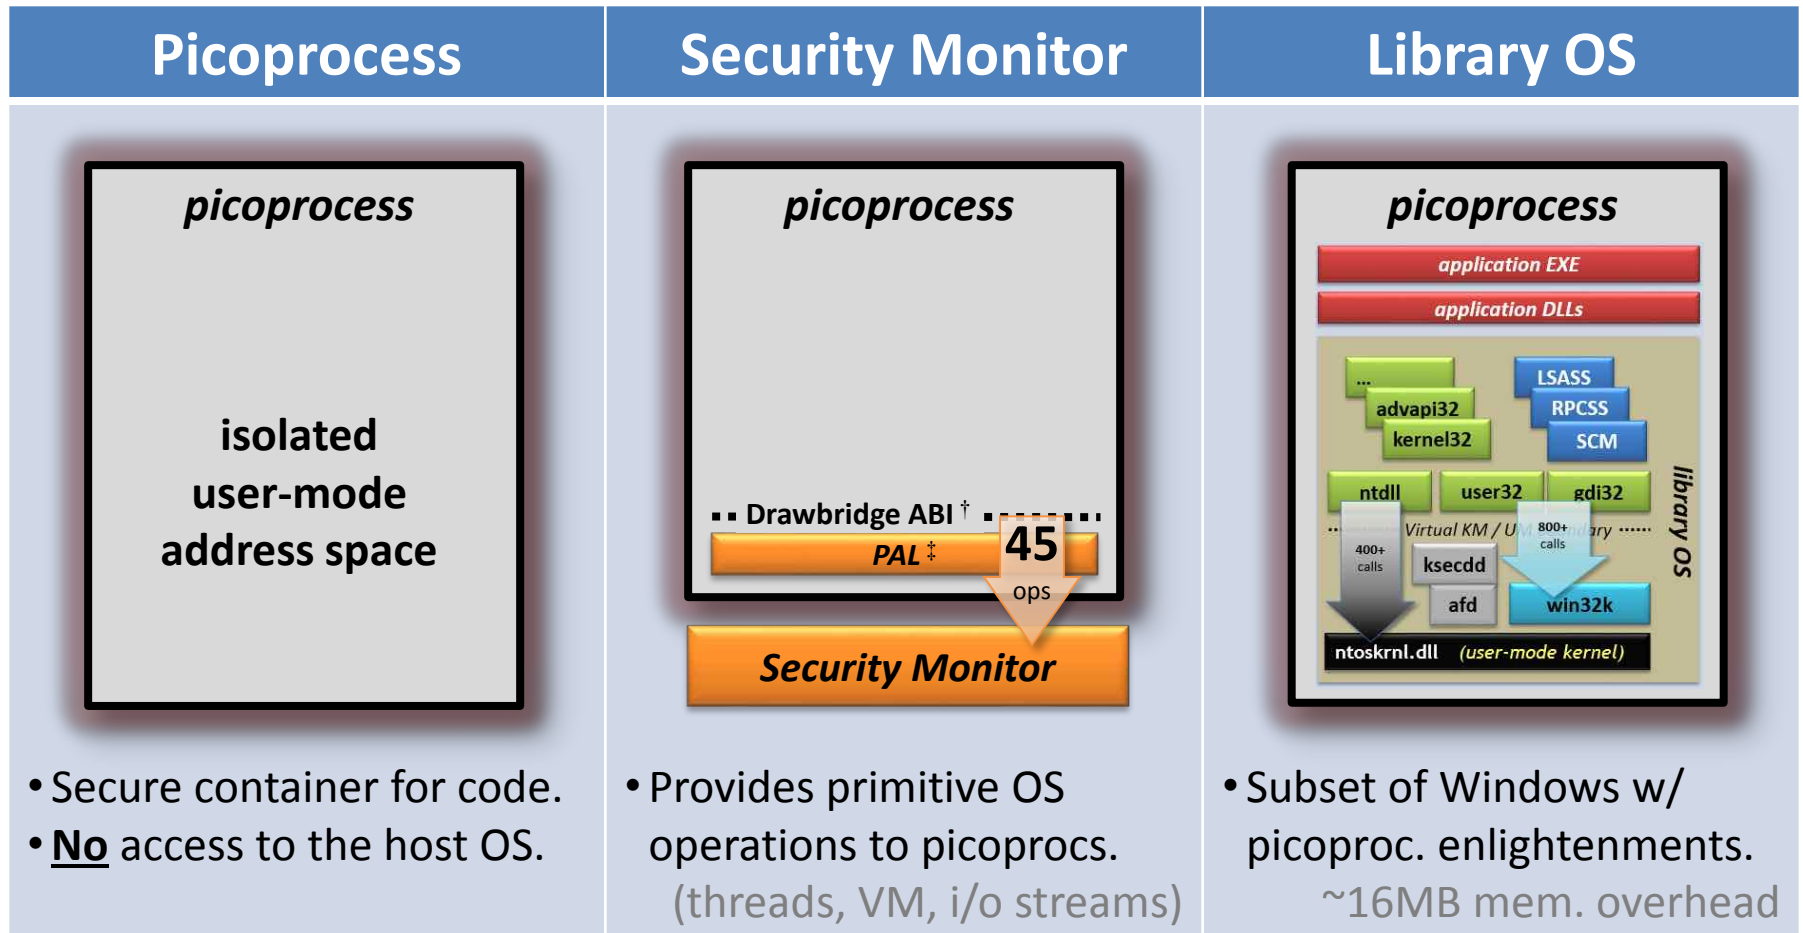
- Adds migration/fault tolerance to any unmodified apps and LibOSes on any Drawbridge host platform
- At runtime, track state:
 - Writable/modified virtual memory allocations
 - All threads & synchronisation
 - Open streams, and their parameters
 - Outstanding I/O
- At checkpoint time:
 - Cancel pending I/O and ABI calls
 - Open file (using ABI)
 - Serialise address space, thread contexts and other state to file

Demo



Summary

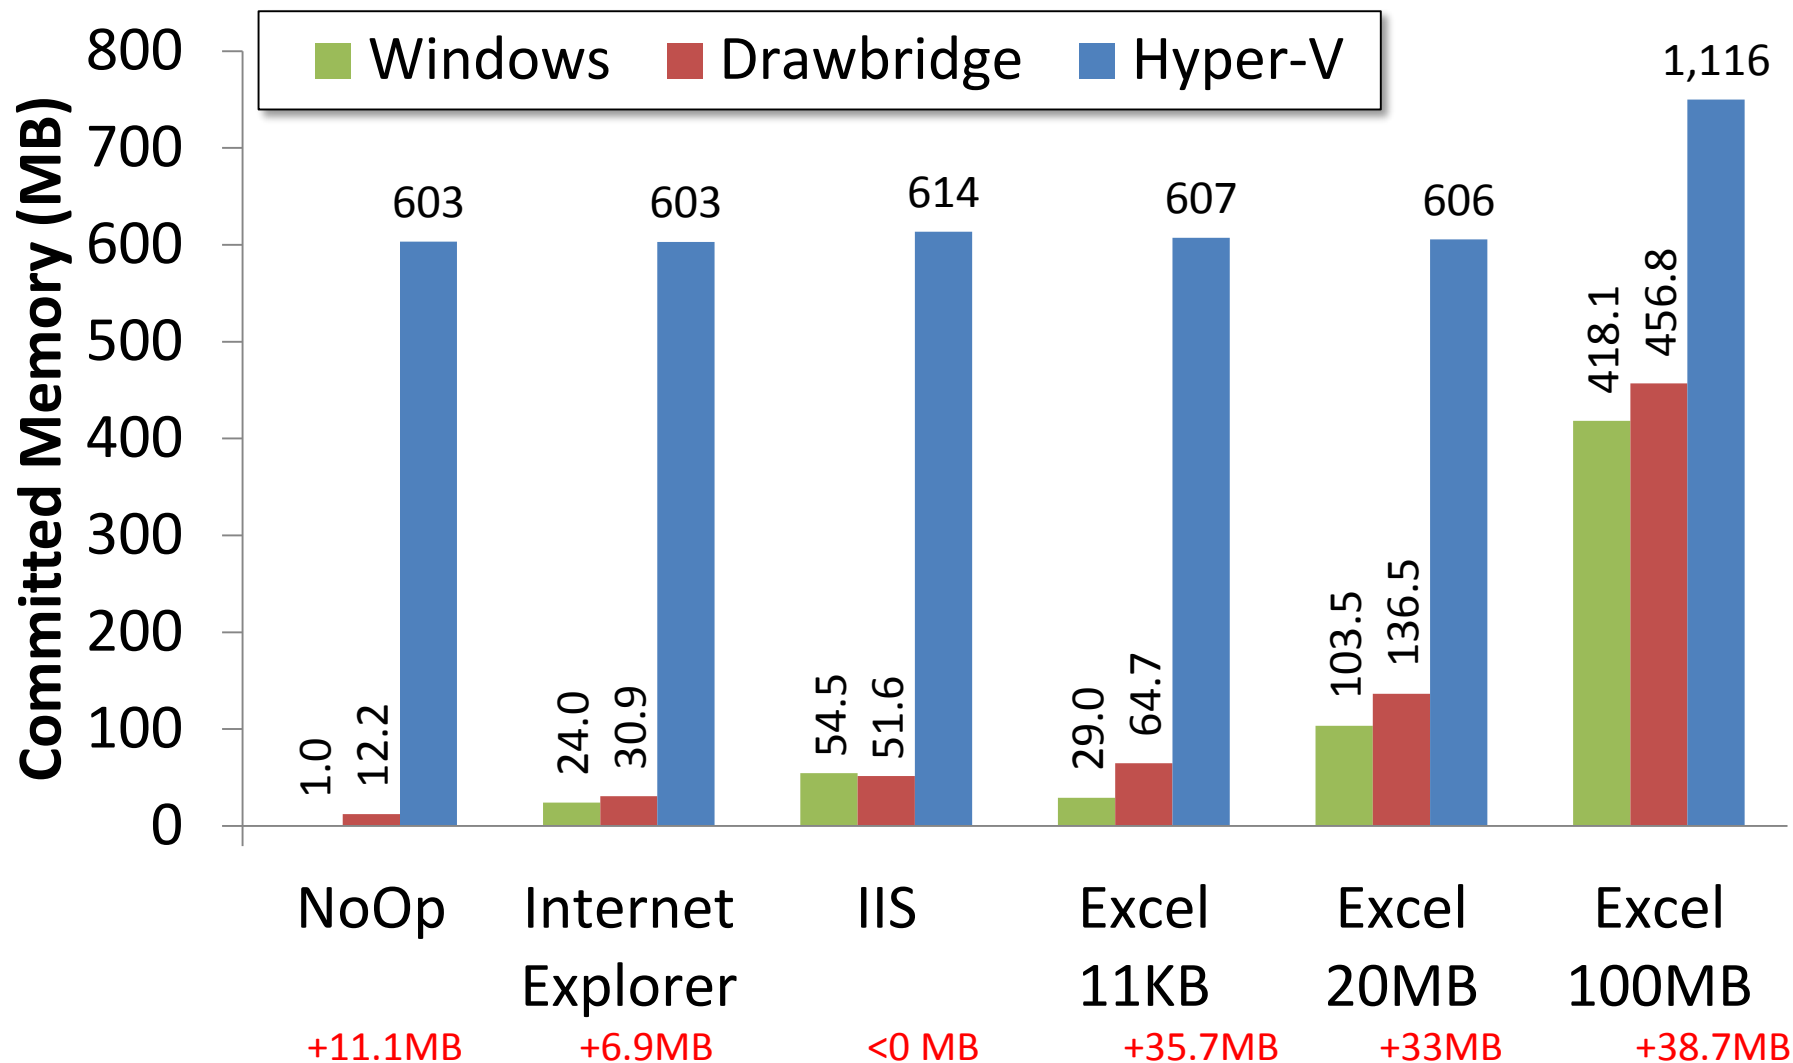
- Drawbridge is a light-weight VMs alternative for secure application hosting.
- Drawbridge consists of three pieces:



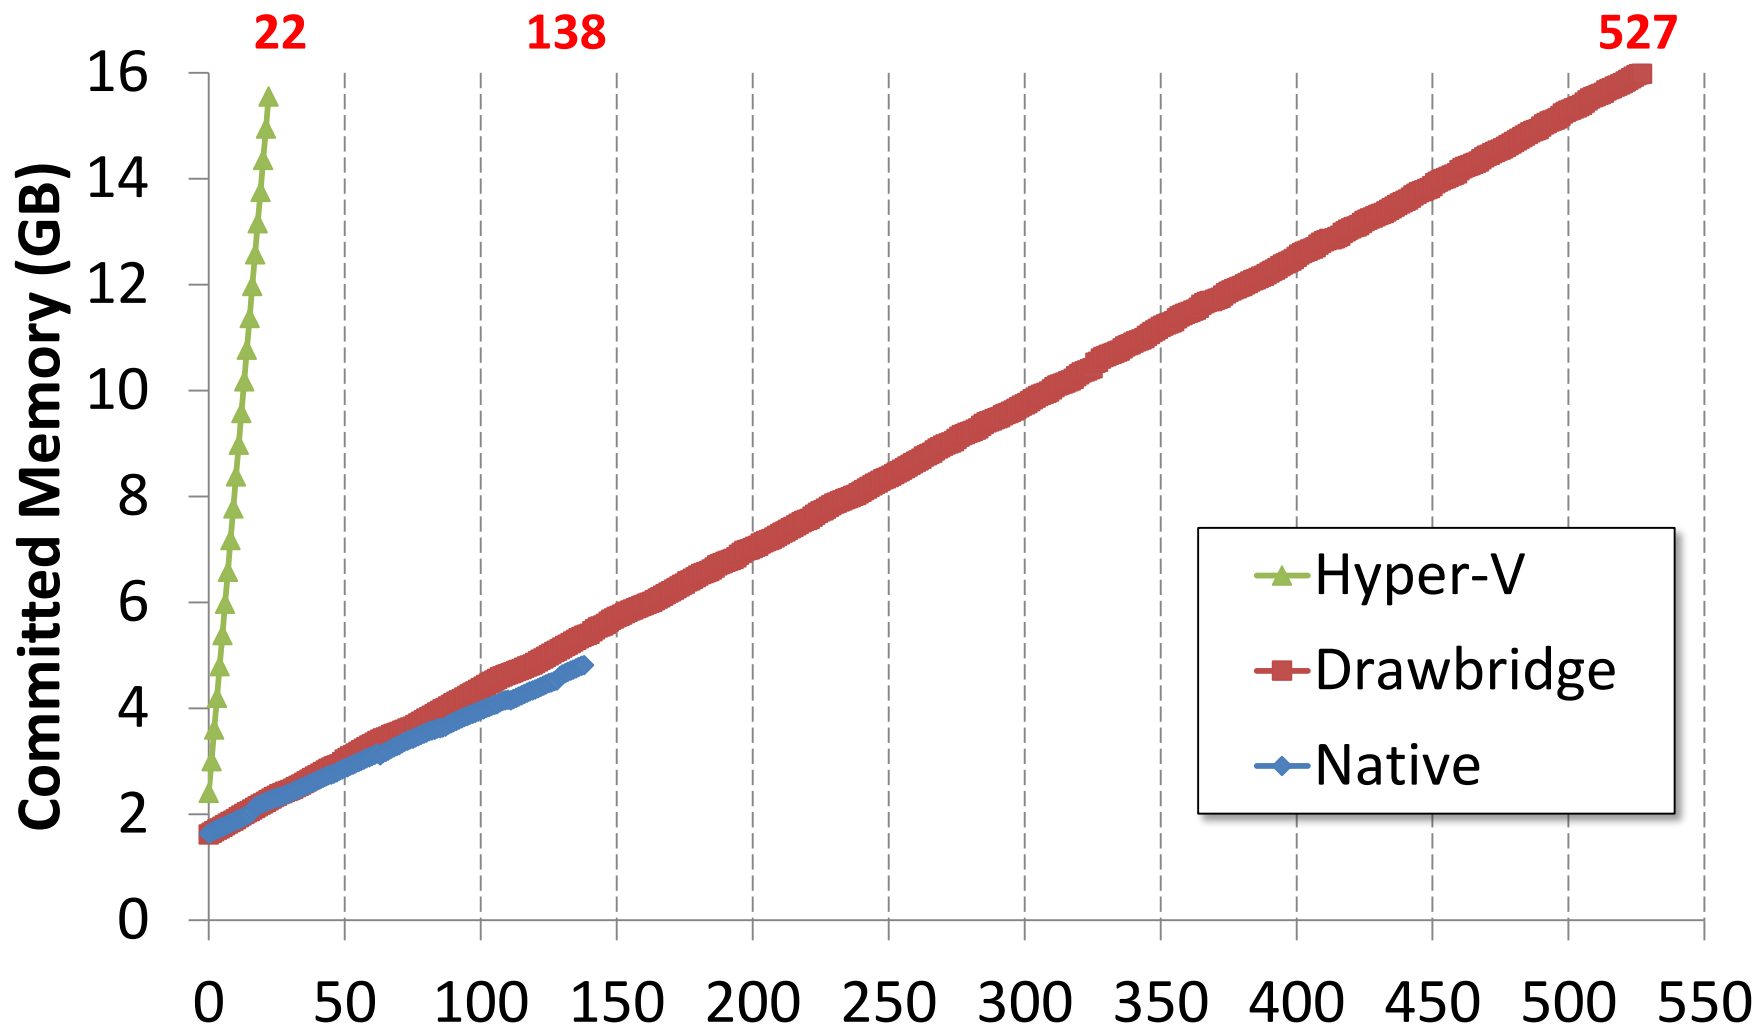
But...

- Do these higher-level abstractions have benefits beyond those of a hardware VM?
- Yes,...
 - Higher density...
 - Layerable (cheaply)...
 - More versatile...
 - ... see also [ASPLOS 2011]

Density: Committed Memory by Application

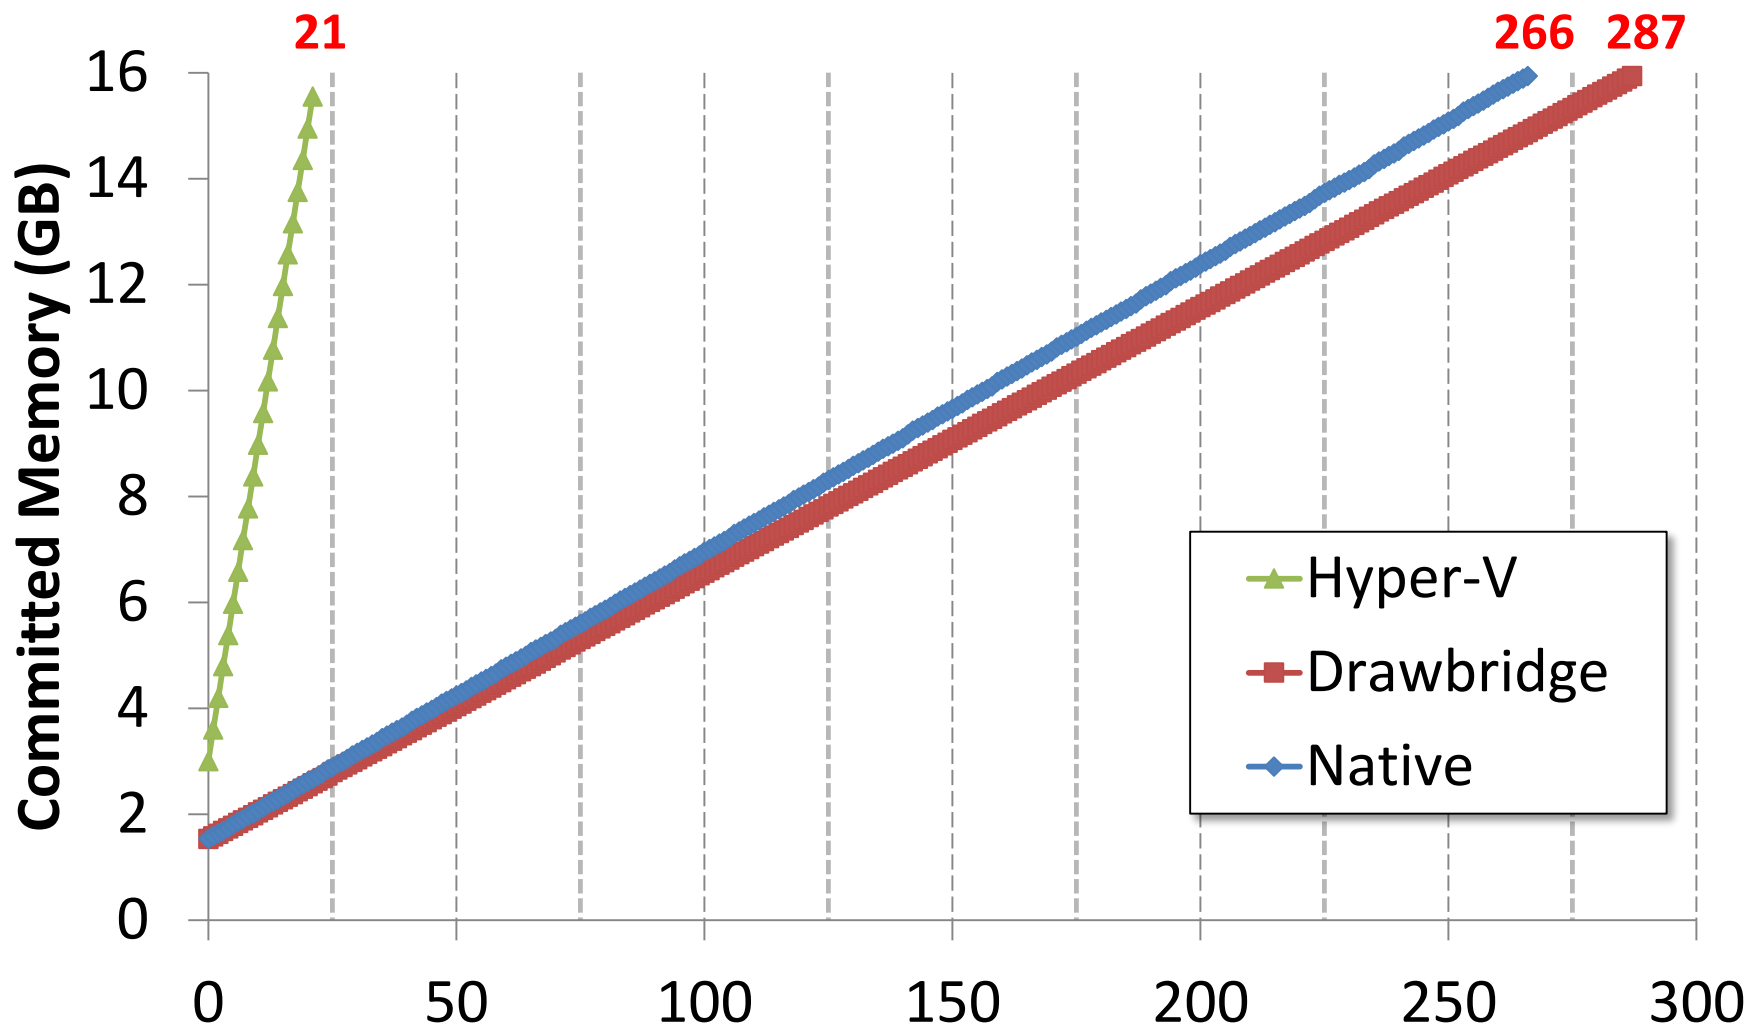


VM Committed for IE8 Instances



* Native stops at 138 instances when IE reaches the per-session limit on GDI handles.

VM Committed for IIS instances

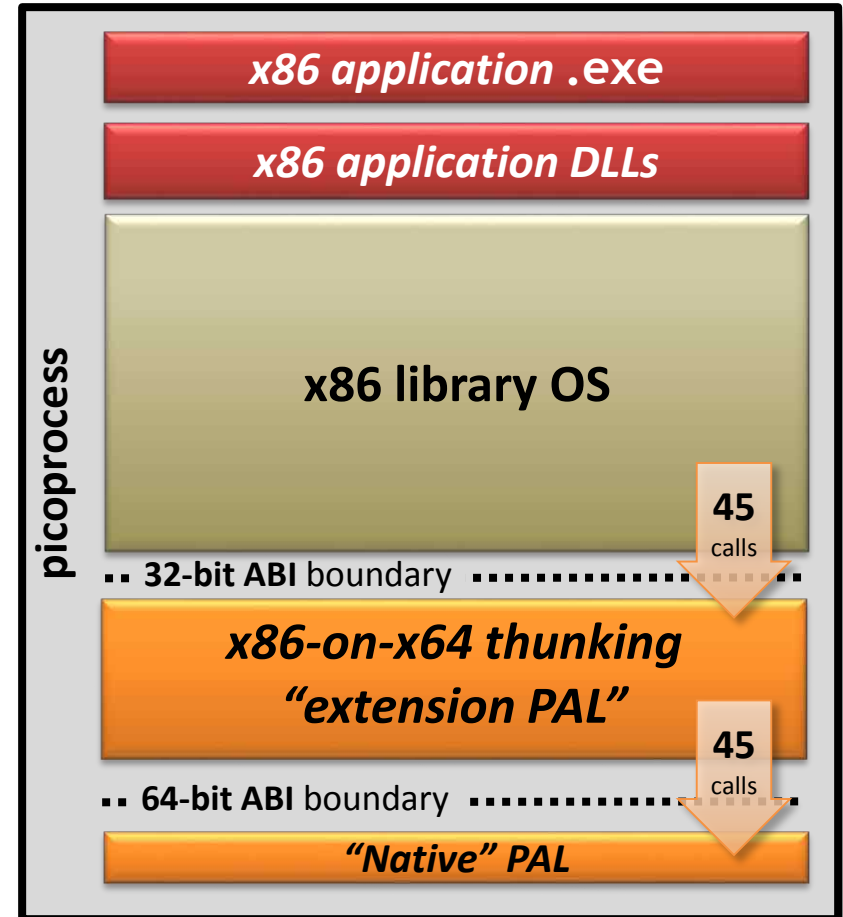


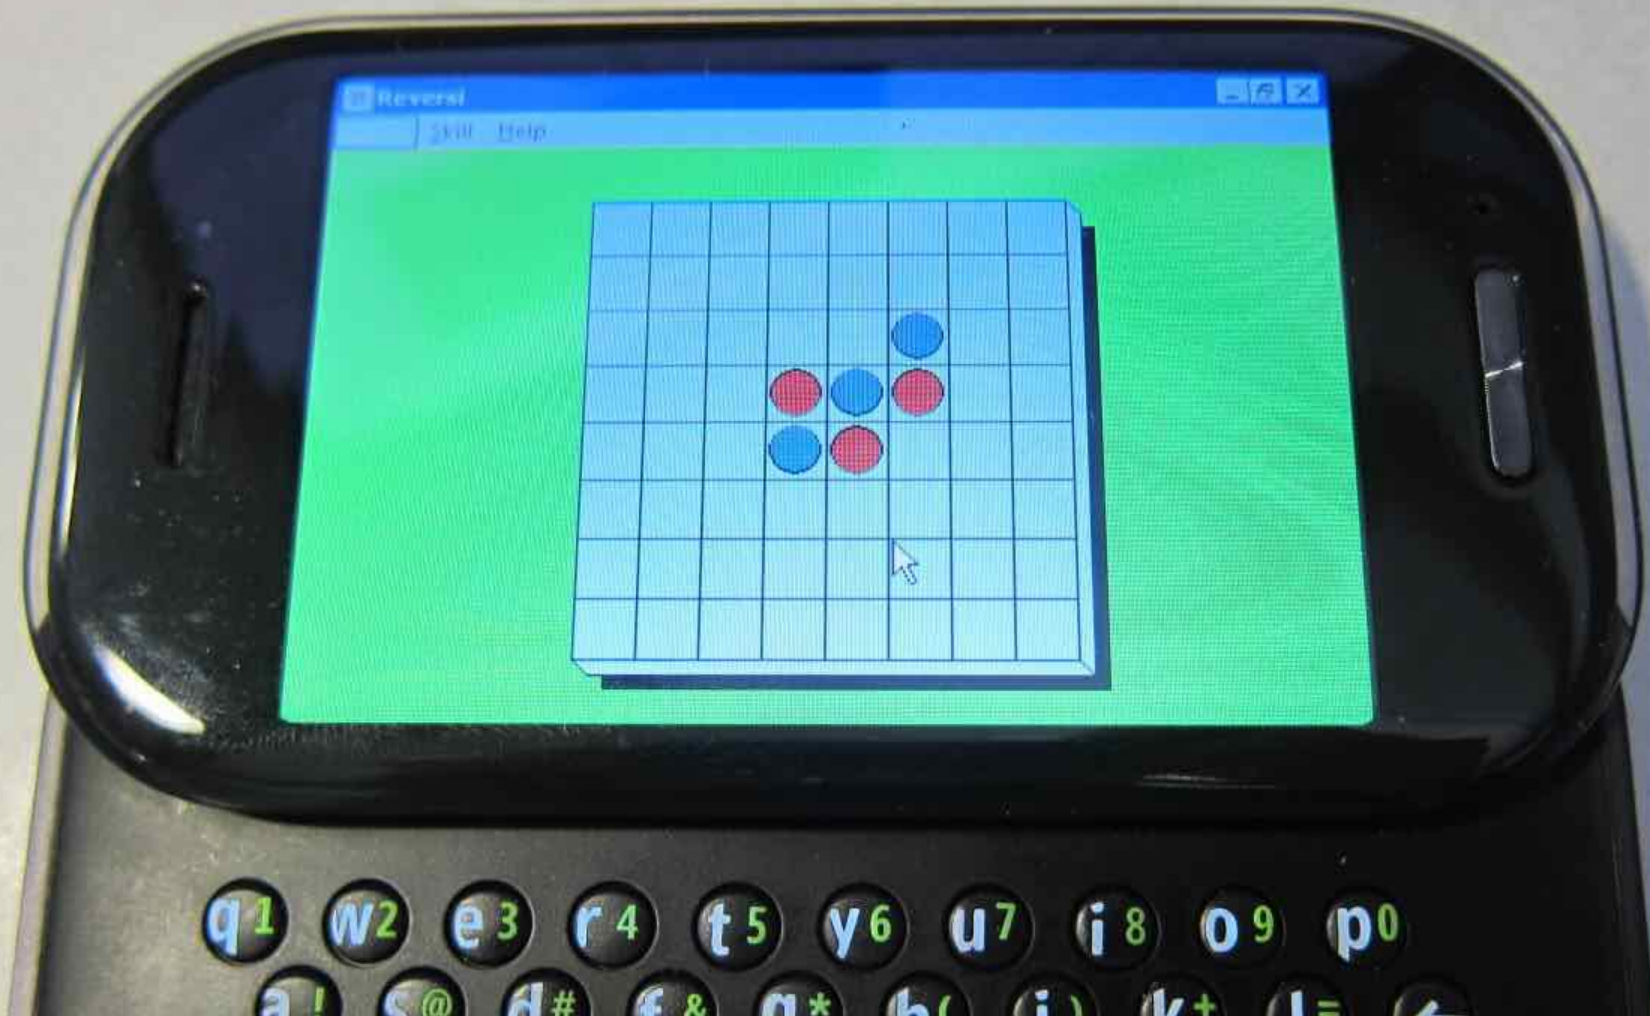
Layerable: Extensions [EuroSys 2013]

- Change the runtime behaviour of an application/OS
- Developed by a third party
- Applied by end user or system integrator
- “Bascule” ABI:
 - Nestable in-process ABI of common OS primitives
 - “Host” provides:
 - Table of function entry points
 - Data structure of startup parameters
 - “Guest” provides:
 - Table of upcall entry points
 - Bootstrap: each layer loads the one above

Drawbridge and WoW64

- Subtly different from WoW64
 - Same 32-bit library OS used on both x86, x64
 - 64-bit host unaware of 32-bit code
- x86-on-x64 “extension” PAL
 - Depends only on 64-bit ABI
 - Exposes 32-bit ABI to library OS
 - Looks like a PAL to the layer above it, like a library OS to the layer below
 - Internally thunks between x86, x64
- Approach generalizes
 - ...for CPU emulation (e.g. x86-on-ARM “emulating extension”)
 - ...for any layered ABI filters





Architecture adaptation extension

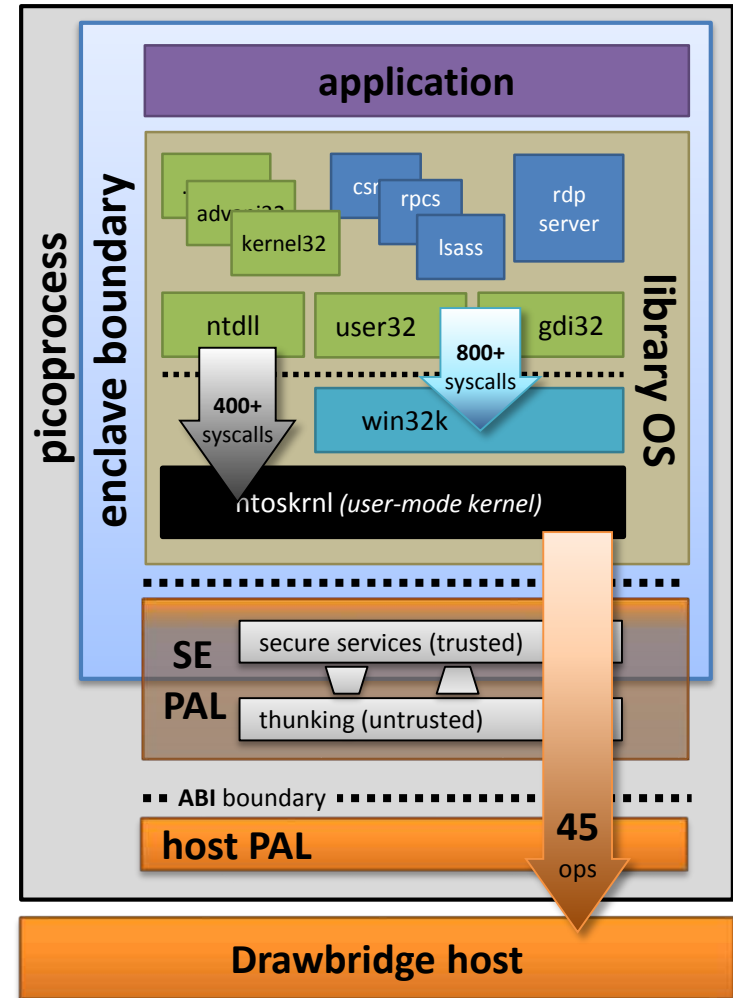
Unmodified x86 app and LibOS on ARM host, migrates from x86 to ARM and back

Sample Extensions [EuroSys 2013]

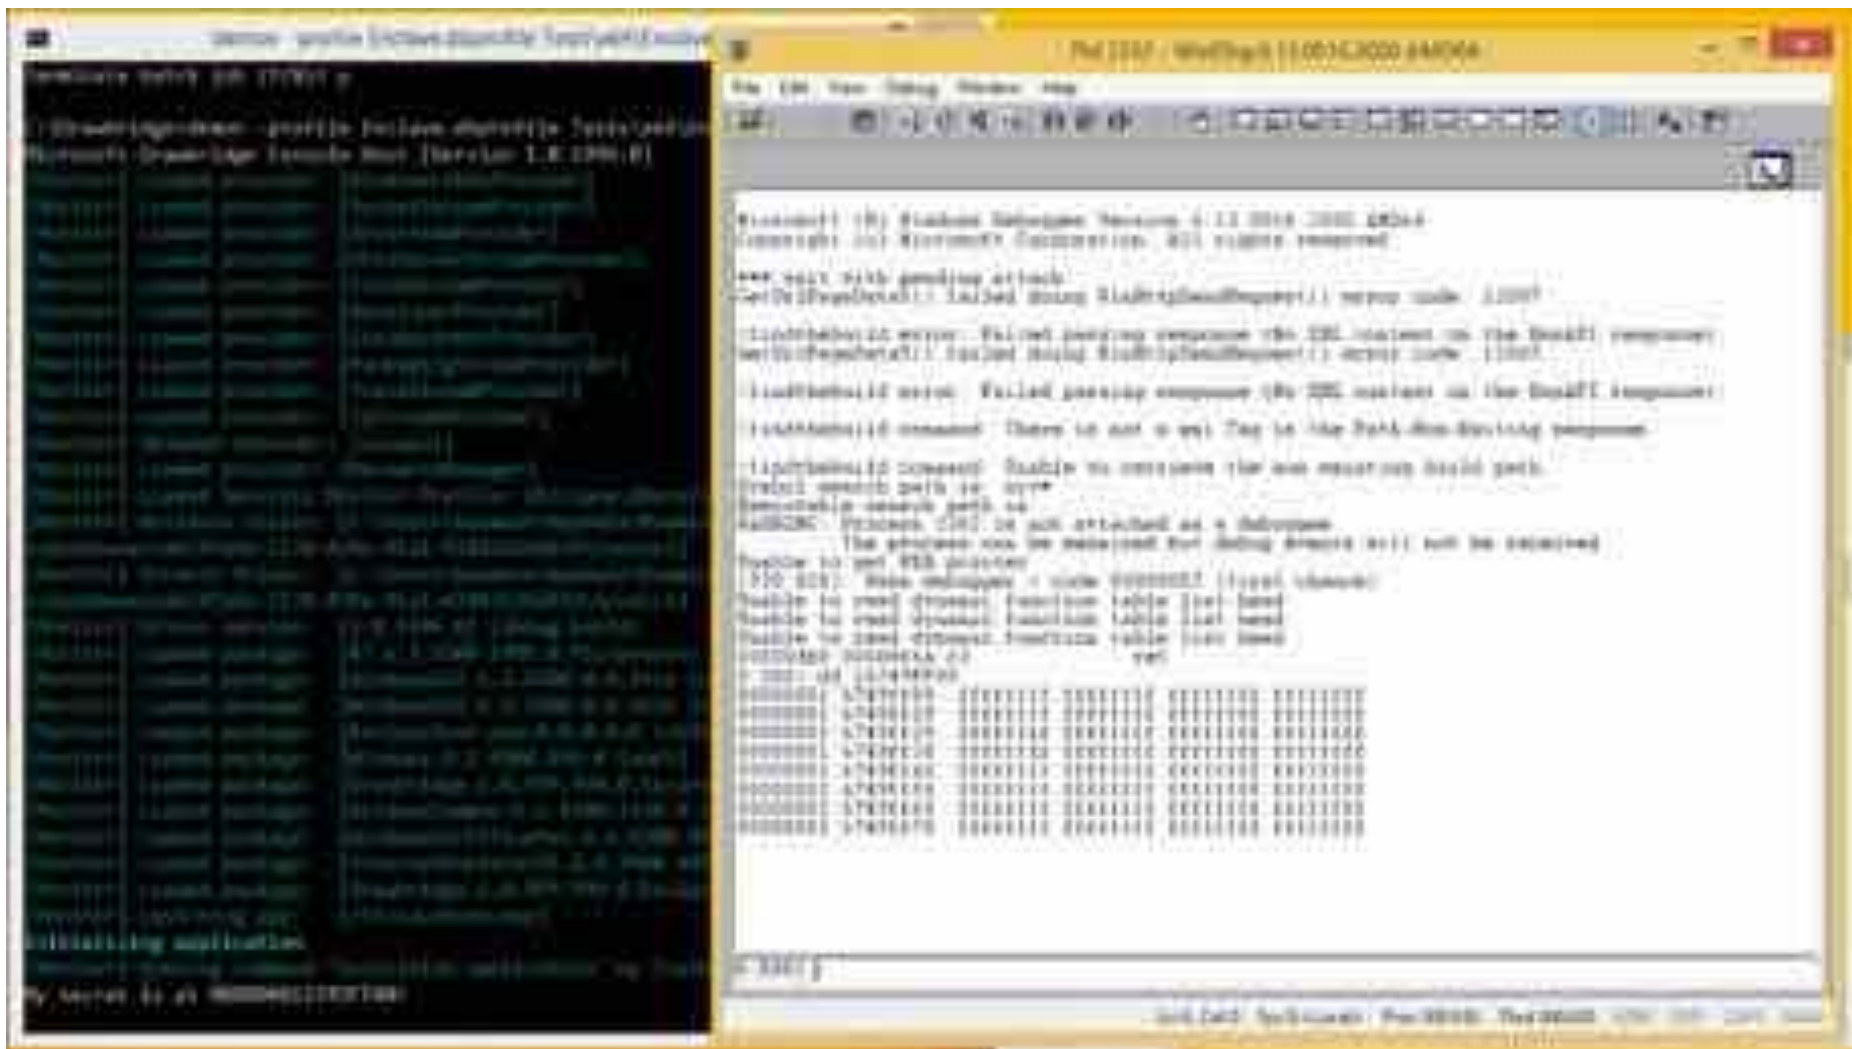
- Implemented:
 - Tracing
 - File system remapping
 - Checkpointing
 - Architecture adaptation
 - 32-on-64-bit x86
 - x86 on ARM JIT
- Discussed:
 - Speculation
 - Record and replay

More Versatile: Intel SGX [SOSP 2013 Demo]

- Problem:
 - Applications must trust:
 - OS, VMM, bootloader, BIOS, etc.
 - Administrator(s), management tools
 - Hierarchical trust model is inadequate
 - No practical way to protect private data
- Intel SGX [HASP 2013]
 - New instructions & memory access changes
 - Confidentiality and integrity protection
 - Protected *enclave* in user-mode process
 - Memory encryption
 - Hardware support for **mutual distrust** in software
 - Remote attestation
 - Secure execution, despite compromised OS/VMM



Intel SGX Demo (Emulator)



The image shows a screenshot of the Intel SGX Demo (Emulator) application. It consists of two main windows side-by-side.

The left window, titled "Secure Binary", displays a list of memory addresses and their corresponding values. The addresses range from 0x00000000 to 0x0000000F. The values are mostly 0x00000000, with some non-zero values at higher addresses, such as 0x00000001, 0x00000002, 0x00000003, 0x00000004, 0x00000005, 0x00000006, 0x00000007, 0x00000008, 0x00000009, 0x0000000A, 0x0000000B, 0x0000000C, 0x0000000D, 0x0000000E, and 0x0000000F.

The right window, titled "Secure Binary (Emulator)", displays a list of memory addresses and their corresponding values. The addresses range from 0x00000000 to 0x0000000F. The values are mostly 0x00000000, with some non-zero values at higher addresses, such as 0x00000001, 0x00000002, 0x00000003, 0x00000004, 0x00000005, 0x00000006, 0x00000007, 0x00000008, 0x00000009, 0x0000000A, 0x0000000B, 0x0000000C, 0x0000000D, 0x0000000E, and 0x0000000F.

Why should we love higher-level abstractions?

- Three reasons:
 - **Compatibility**
 - I can install my application in a picoprocess with the Library OS it requires and never have to worry about it breaking again.
 - **Security**
 - I can download even the most malignant code from the internet, run it in a picoprocess, and my system's integrity isn't lost.
 - **Continuity**
 - I can start an application (in a picoprocess) on one computer and move it to another, or reboot the computer, and it still runs.
- And three more:
 - **High Density**
 - I can run many hundreds of picoprocesses on one computer.
 - **Layerable**
 - I can write an extension that works with many hosts or library Oses.
 - **Versatility**
 - I can use higher-level abstractions where a VM can't run (cheaply).

