

**ISA 673**  
**Operating Systems' Security**

# Introduction to the Pin Instrumentation Tool

Quan Jia

Mar 27, 2013

# What is Pin?

*Pin is Intel's dynamic binary instrumentation engine.*

# What is Instrumentation?

- A technique that inserts extra code into a program to collect runtime information.
- Program analysis : performance profiling, error detection, capture & replay
- Architectural study : processor and cache simulation, trace collection
- Binary translation : Modify program behavior, emulate unsupported instructions


# Instrumentation Approaches

- Source Code Instrumentation (SCI)
  - instrument source programs
- Binary Instrumentation (BI)
  - instrument binary executable directly

# SCI Example (Code Coverage)

## Original Program

```
void foo() {  
    bool found=false;  
    for (int i=0; i<100; ++i) {  
        if (i==50) break;  
        if (i==20) found=true;  
    }  
    printf("foo\n");  
}
```



## Instrumented Program

```
char inst[5];  
void foo() {  
    bool found=false; inst[0]=1;  
    for (int i=0; i<100; ++i) {  
        if (i==50) { inst[1]=1;break;}  
        if (i==20) { inst[2]=1;found=true;}  
        inst[3]=1;  
    }  
    printf("foo\n");  
    inst[4]=1;  
}
```

# Binary Instrumentation (BI)

- Static binary instrumentation – inserts additional code and data before execution and generates a persistent modified executable
- Dynamic binary instrumentation – inserts additional code and data during execution without making any permanent modifications to the executable.

# BI Example – Instruction Count

```
counter++;  
sub $0xff, %edx  
counter++;  
cmp %esi, %edx  
counter++;  
jle <L1>  
counter++;  
mov $0x1, %edi  
counter++;  
add $0x10, %eax
```

# BI Example – Instruction Trace

```
Print(ip);  
sub $0xff, %edx  
Print(ip);  
cmp %esi, %edx  
Print(ip);  
jle <L1>  
Print(ip);  
mov $0x1, %edi  
Print(ip);  
add $0x10, %eax
```



# Advantages

- Binary instrumentation
  - Language independent
  - Machine-level view
  - Instrument legacy/proprietary software
- Dynamic instrumentation
  - No need to recompile or relink
  - Discover code at runtime
  - Handle dynamically-generated code
  - Attach to running processes

# What is Pin?

*Pin is Intel's dynamic binary instrumentation engine.*

# Advantages of Pin Instrumentation

- **Easy-to-use Instrumentation:**
  - Uses dynamic instrumentation - Do not need source code, recompilation, post-linking
- **Programmable Instrumentation:**
  - Provides rich APIs to write in C/C++ your own instrumentation tools (called Pintools)
- **Multiplatform:**
  - Supports x86, x86-64, Itanium, Xscale
  - OS's: Windows, Linux, OSX, Android
- **Robust:**
  - Instruments real-life applications: Database, web browsers, ...
  - Instruments multithreaded applications
  - Supports signals
- **Efficient:**
  - Applies compiler optimizations on instrumentation code

# Widely Used and Supported

- Large user base in academia and industry
  - 30,000+ downloads
  - 700+ citations
  - Active mailing list (Pinheads)
- Actively developed at Intel
  - Intel products and internal tools depend on it
  - Nightly testing of 25000 binaries on 15 platforms

# Using Pin

Launch and instrument an application

```
$ pin -t pintool.so -- application
```

↑  
Instrumentation engine  
(provided in the kit)

↖  
Instrumentation tool  
(write your own, or use one  
provided in the kit)

Attach to and instrument an application

```
$ pin -t pintool.so -pid 1234
```

# Pin and Pintools

- Pin – the instrumentation **engine**
- Pintool – the instrumentation **program**
- Pin provides the framework and API, Pintools run on Pin to perform meaningful tasks.
- Pintools
  - Written in C/C++ using Pin APIs
  - Many open source examples provided with the Pin kit
  - Certain Do's and Don'ts apply

# Pin Instrumentation Capabilities

- Replace application functions with your own.
- Fully examine any application instruction – insert a call to your instrumenting function whenever that instruction executes.
- Pass a large set of supported parameters to your instrumenting function.
  - Register values (including IP), Register values by reference (for modification)
  - Memory addresses read/written by the instruction
  - Full register context
- Track function calls including syscalls and examine/change arguments.
- Track application threads.
- Intercept signals.
- Instrument a process tree.

# Hands-on Task

- Download the latest Pin from <http://www.pintool.org>
  - For Windows: make sure you download the correct version that matches your Visual Studio IDE.
- Build all included Pintools under
  - [source/tools/SimpleExamples](#)
  - [source/tools/ManualExmaples](#)
- Refer to the user's manual for detailed instructions
  - **Attention:** Nmake does not work for Windows, use Cygwin to install GNU make instead.



# Pintool 1: Instruction Count

```
counter++;  
sub $0xff, %edx  
counter++;  
cmp %esi, %edx  
counter++;  
jle <L1>  
counter++;  
mov $0x1, %edi  
counter++;  
add $0x10, %eax
```

# Pintool 1: Invocation

- Windows examples:

```
> pin.exe -t inscount0.dll -- dir.exe
```

```
> pin.exe -t inscount0.dll -o incount.out -- gzip.exe FILE
```

- Linux examples:

```
$ pin -t inscount0.so -- /bin/ls
```

```
$ pin -t inscount0.so -o incount.out -- gzip FILE
```

# Pintool 1: ManualExamples/inscount0.cpp

```
#include <iostream>
#include "pin.h"
```

```
UINT64 icount = 0;
```

```
void docount() { icount++; }
```

*analysis routine*

```
void Instruction(INS ins, void *v)
```

*instrumentation routine*

```
{
    INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)docount, IARG_END);
}
```

```
void Fini(INT32 code, void *v)
```

```
{ std::cerr << "Count " << icount << endl; }
```

switch to pin stack  
save registers  
call docount  
restore registers  
switch to app stack

```
int main(int argc, char * argv[])
```

```
{
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

# Pin Instrumentation APIs

- Basic APIs are architecture independent:
  - Provide common functionalities like determining:
    - Control-flow changes
    - Memory accesses
- Architecture-specific APIs
  - E.g., Info about segmentation registers on IA32
- Call-based APIs:
  - Instrumentation routines
  - Analysis routines

# Pintool 2: Instruction Trace

```
Print(ip);  
sub $0xff, %edx  
Print(ip);  
cmp %esi, %edx  
Print(ip);  
jle <L1>  
Print(ip);  
mov $0x1, %edi  
Print(ip);  
add $0x10, %eax
```

# Pintool 2:

[ManualExamples/itrace.cpp](#)

```
#include <stdio.h>
#include "pin.H"
FILE * trace;
```

```
void printip(void *ip) { fprintf(trace, "%p\n", ip); }
```

argument to analysis routine

*analysis routine*

```
void Instruction(INS ins, void *v) {
    INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)printip,
                  IARG_INST_PTR, IARG_END);
}
```

*instrumentation routine*

```
void Fini(INT32 code, void *v) { fclose(trace); }

int main(int argc, char * argv[]) {
    trace = fopen("itrace.out", "w");
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);

    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

# Examples of Arguments to Analysis Routine

IARG\_INST\_PTR

- **Instruction pointer (program counter) value**

IARG\_UINT32 <value>

- **An integer value**

IARG\_REG\_VALUE <register name>

- **Value of the register specified**

IARG\_BRANCH\_TARGET\_ADDR

- **Target address of the branch instrumented**

IARG\_MEMORY\_READ\_EA


- **Effective address of a memory read**

*And many more ... (refer to the Pin manual for details)*

# Instrumentation Points

Instrument points relative to an instruction:

- *Before (IPOINT\_BEFORE)*
- After:
  - Fall-through edge (IPOINT\_AFTER)
  - **Taken edge (IPOINT\_TAKEN)**





# Instrumentation Granularity

**Instrumentation can be done at three different granularities:**

- **Instruction**
- **Basic block**
  - A sequence of instructions terminated at a control-flow changing instruction
  - Single entry, single exit
- **Trace**
  - A sequence of basic blocks terminated at an unconditional control-flow changing instruction
  - Single entry, multiple exits

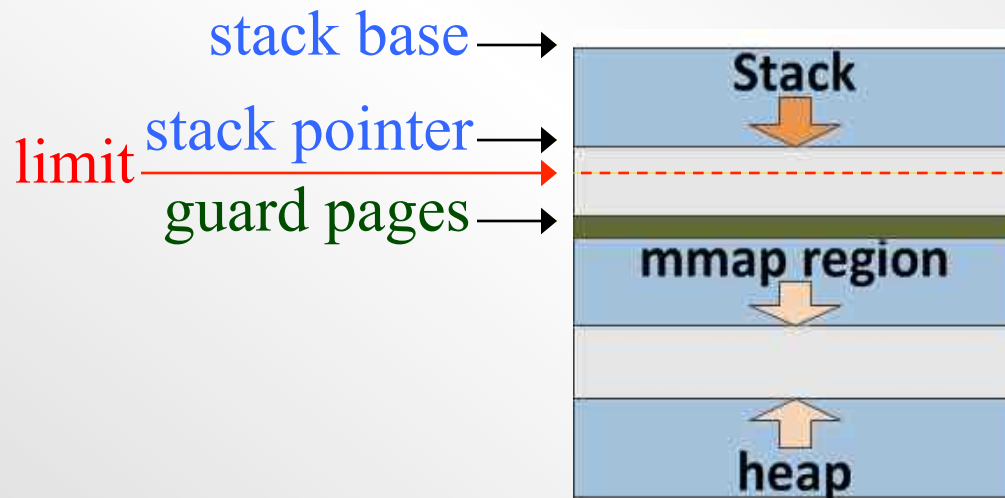
```
sub    $0xff, %edx  
cmp    %esi, %edx  
jle    <L1>
```

```
mov    $0x1, %edi  
add    $0x10, %eax  
jmp    <L2>
```

1 Trace, 2 BBs, 6 insts

# Hands-on Task: Stack Monitor

- **Goal:** Monitor runtime stack usage and alert if it exceeds a pre-defined limit.
- Process address space:



# Hands-on Task: Stack Monitor

- **Steps:**
  1. Obtain stack base address when process starts.
  2. Perform instruction-level instrumentation.
  3. Get runtime stack size (`stack_base` – `stack_pointer`).
  4. Compare stack size with supplied size limit.
- **Hint:** refer to `ManualExamples/stack-debugger.cpp`