

# Understanding the Python GIL

David Beazley  
<http://www.dabeaz.com>

Presented at PyCon 2010  
Atlanta, Georgia

# Introduction

- As a few of you might know, C Python has a Global Interpreter Lock (GIL)

```
>>> import that
```

```
The Unwritten Rules of Python
```

1. You do not talk about the GIL.
2. You do NOT talk about the GIL.
3. Don't even mention the GIL. No seriously.
- ...

- It limits thread performance
- Thus, a source of occasional "contention"

# An Experiment

- Consider this trivial CPU-bound function

```
def countdown(n):  
    while n > 0:  
        n -= 1
```

- Run it once with a lot of work

```
COUNT = 100000000    # 100 million  
countdown(COUNT)
```

- Now, subdivide the work across two threads

```
t1 = Thread(target=countdown, args=(COUNT//2,))  
t2 = Thread(target=countdown, args=(COUNT//2,))  
t1.start(); t2.start()  
t1.join(); t2.join()
```

# A Mystery

- Performance on a quad-core MacPro

Sequential : 7.8s

Threaded (2 threads) : 15.4s (2X slower!)

- Performance if work divided across 4 threads

Threaded (4 threads) : 15.7s (about the same)

- Performance if all but one CPU is disabled

Threaded (2 threads) : 11.3s (~35% faster than running

Threaded (4 threads) : 11.6s with all 4 cores)

- Think about it...

# This Talk

- An in-depth look at threads and the GIL that will explain that mystery and much more
- Some cool pictures
- A look at the new GIL in Python 3.2

# Disclaimers

- I gave an earlier talk on this topic at the Chicago Python Users Group (chipy)  
<http://www.dabeaz.com/python/GIL.pdf>
- That is a different, but related talk
- I'm going to go pretty fast... please hang on

# Part I

## Threads and the GIL

# Python Threads

- Python threads are real system threads
  - POSIX threads (pthreads)
  - Windows threads
- Fully managed by the host operating system
- Represent threaded execution of the Python interpreter process (written in C)




# Alas, the GIL

- Parallel execution is forbidden
- There is a "global interpreter lock"
- The GIL ensures that only one thread runs in the interpreter at once
- Simplifies many low-level details (memory management, callouts to C extensions, etc.)

# Thread Execution Model

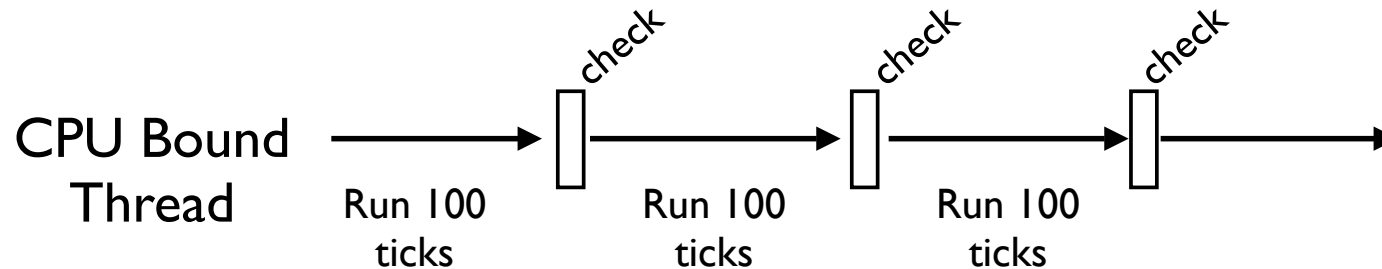
- With the GIL, you get cooperative multitasking



- When a thread is running, it holds the GIL
- GIL released on I/O (read,write,send,recv,etc.)

# CPU Bound Tasks

- CPU-bound threads that never perform I/O are handled as a special case
- A "check" occurs every 100 "ticks"



- Change it using `sys.setcheckinterval()`

# What is a "Tick?"

- Ticks loosely map to interpreter instructions

```
def countdown(n):  
    while n > 0:  
        print n  
        n -= 1
```

- Instructions in the Python VM
- Not related to timing (ticks might be long)

	>>> <b>import dis</b>	
	>>> <b>dis.dis(countdown)</b>	
	0 SETUP_LOOP	33 (to 36)
	3 LOAD_FAST	0 (n)
	6 LOAD_CONST	1 (0)
	9 COMPARE_OP	4 (>)
Tick 1	12 JUMP_IF_FALSE	19 (to 34)
	15 POP_TOP	
	16 LOAD_FAST	0 (n)
	19 PRINT_ITEM	
Tick 2	20 PRINT_NEWLINE	
	21 LOAD_FAST	0 (n)
Tick 3	24 LOAD_CONST	2 (1)
	27 INPLACE_SUBTRACT	
	28 STORE_FAST	0 (n)
Tick 4	31 JUMP_ABSOLUTE	3
	...	

# The Periodic "Check"

- The periodic check is really simple
- The currently running thread...
  - Resets the tick counter
  - Runs signal handlers if the main thread
  - Releases the GIL
  - Reacquires the GIL
- That's it

# Implementation (C)

Note: Each thread is running this same code

```
/* Python/ceval.c */
...
Decrement ticks    if (--_Py_Ticker < 0) {
                    ...
Reset ticks        _Py_Ticker = _Py_CheckInterval;
                    ...
Run signal handlers if (things_to_do) {
                    if (Py_MakePendingCalls() < 0) {
                        ...
                    }
                    }
Release and reacquire the GIL if (interpreter_lock) {
                            /* Give another thread a chance */
                            PyThread_release_lock(interpreter_lock);

                            /* Other threads may run now */

                            PyThread_acquire_lock(interpreter_lock, 1);
                    }
                    ...
```

# Big Question

- What is the source of that large CPU-bound thread performance penalty?
- There's just not much code to look at
- Is GIL acquire/release solely responsible?
- How would you find out?

# Part 2

## The GIL and Thread Switching Deconstructed



# Python Locks

- The Python interpreter only provides a single lock type (in C) that is used to build all other thread synchronization primitives
- It's not a simple mutex lock
- It's a binary semaphore constructed from a pthreads mutex and a condition variable
- The GIL is an instance of this lock

# Locks Deconstructed

- Locks consist of three parts

```
locked = 0                # Lock status
mutex  = pthreads_mutex() # Lock for the status
cond   = pthreads_cond()  # Used for waiting/wakeup
```

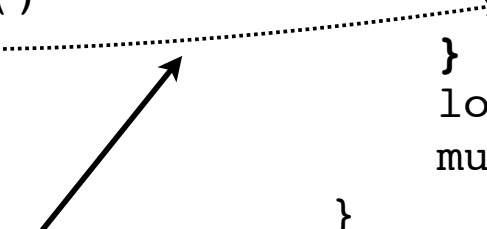
- Here's how acquire() and release() work

pseudocode

```
release() {
    mutex.acquire()
    locked = 0
    mutex.release()
    cond.signal()
}
```

```
acquire() {
    mutex.acquire()
    while (locked) {
        cond.wait(mutex)
    }
    locked = 1
    mutex.release()
}
```

A critical aspect  
concerns this signaling  
between threads



# Thread Switching

- Suppose you have two threads


Thread 1  $\xrightarrow{\text{Running}}$

Thread 2 READY

- Thread 1 : Running
- Thread 2 : Ready (Waiting for GIL)

# Thread Switching

- Easy case : Thread 1 performs I/O (read/write)

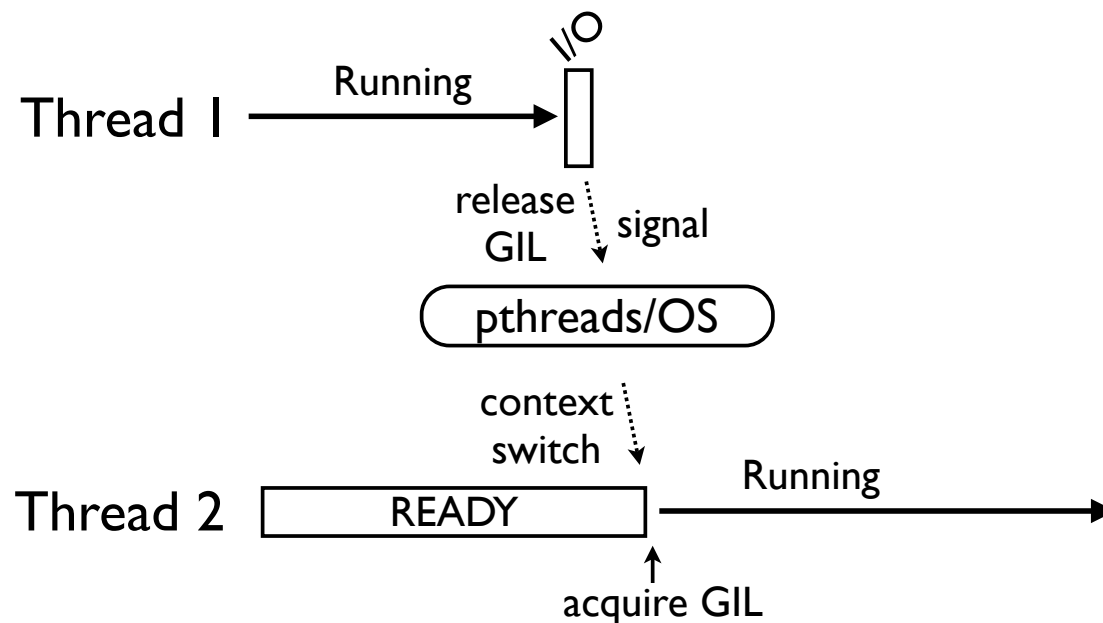


Thread 2 READY

- Thread 1 might block so it releases the GIL

# Thread Switching


- Easy case : Thread 1 performs I/O (read/write)



- Release of GIL results in a signaling operation
- Handled by thread library and operating system

# Thread Switching

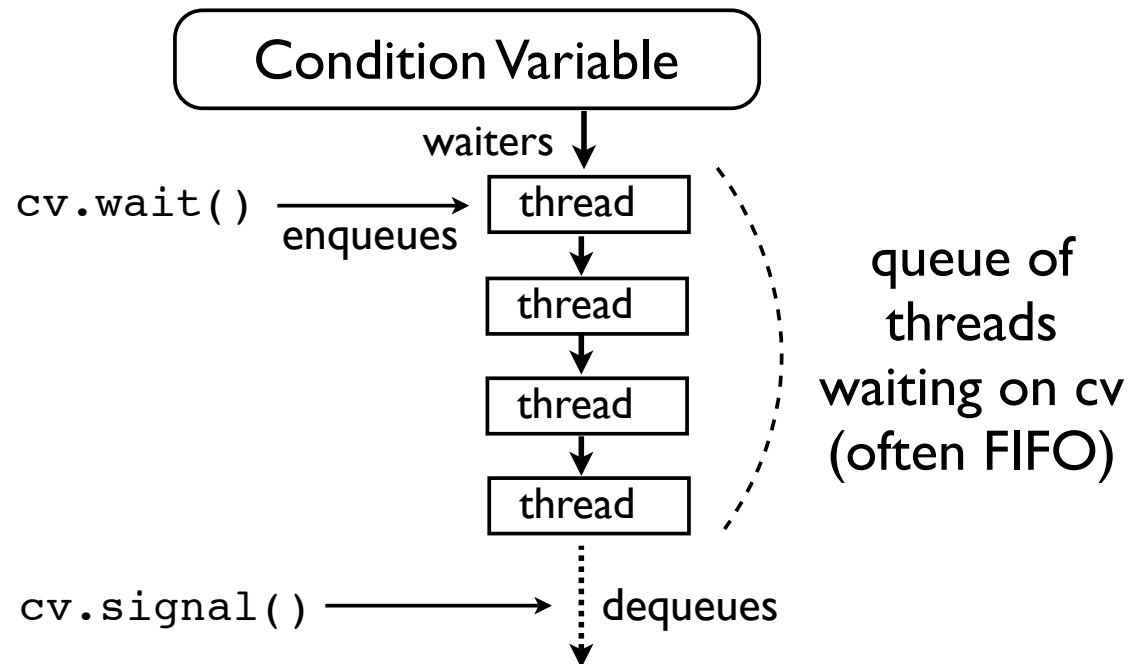
- Tricky case : Thread 1 runs until the check



- Either thread is able to run
- So, which is it?

# pthread Undercover

- Condition variables have an internal wait queue



- Signaling pops a thread off of the queue
- However, what happens after that?

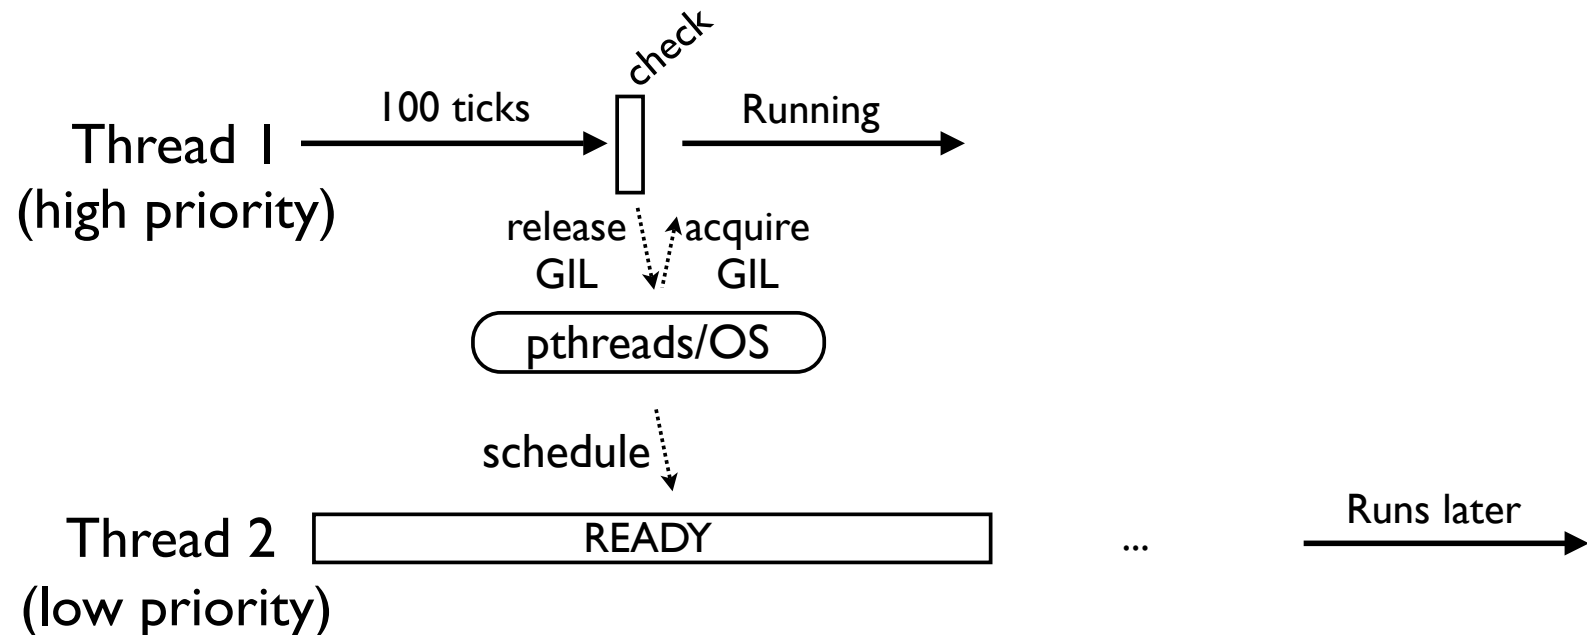
# OS Scheduling

- The operating system has a priority queue of threads/processes ready to run
- Signaled threads simply enter that queue
- The operating system then runs the process or thread with the highest priority
- It may or may not be the signaled thread



# Thread Switching


- Thread 1 might keep going



- Thread 2 moves to the OS "ready" queue and executes at some later time

# Thread Switching

- Thread 2 might immediately take over



- Again, highest priority wins

# Part 3

## What Can Go Wrong?

# GIL Instrumentation

- To study thread scheduling in more detail, I instrumented Python with some logging
- Recorded a large trace of all GIL acquisitions, releases, conflicts, retries, etc.
- Goal was to get a better idea of how threads were scheduled, interactions between threads, internal GIL behavior, etc.

# GIL Logging

- An extra tick counter was added to record number of cycles of the check interval
- Locks modified to log GIL events (pseudocode)

```
release() {  
    mutex.acquire()  
    locked = 0  
    if gil: log("RELEASE")  
    mutex.release()  
    cv.signal()  
}
```

Note: Actual code in C, event logs are stored entirely in memory until exit (no I/O)

```
acquire() {  
    mutex.acquire()  
    if locked and gil:  
        log("BUSY")  
    while locked:  
        cv.wait(mutex)  
        if locked and gil:  
            log("RETRY")  
    locked = 1  
    if gil: log("ACQUIRE")  
    mutex.release()  
}
```

# A Sample Trace

thread id →	t2	100	5351	<b>ACQUIRE</b>	← ACQUIRE : GIL acquired
	t2	100	5352	<b>RELEASE</b>	← RELEASE : GIL released
	t2	100	5352	ACQUIRE	
	t2	100	5353	RELEASE	
	t1	100	5353	ACQUIRE	
tick	t2	<b>38</b>	5353	<b>BUSY</b>	← BUSY : Attempted to acquire
countdown	t1	100	5354	RELEASE	GIL, but it was already in use
	t1	100	5354	ACQUIRE	
	t2	79	5354	RETRY	
	t1	100	5355	RELEASE	
total	t1	100	<b>5355</b>	ACQUIRE	
number of	t2	73	5355	<b>RETRY</b>	← RETRY : Repeated attempt to
"checks"	t1	100	5356	RELEASE	acquire the GIL, but it was
executed	t2	100	5356	ACQUIRE	still in use
	t1	24	5356	BUSY	
	t2	100	5357	RELEASE	

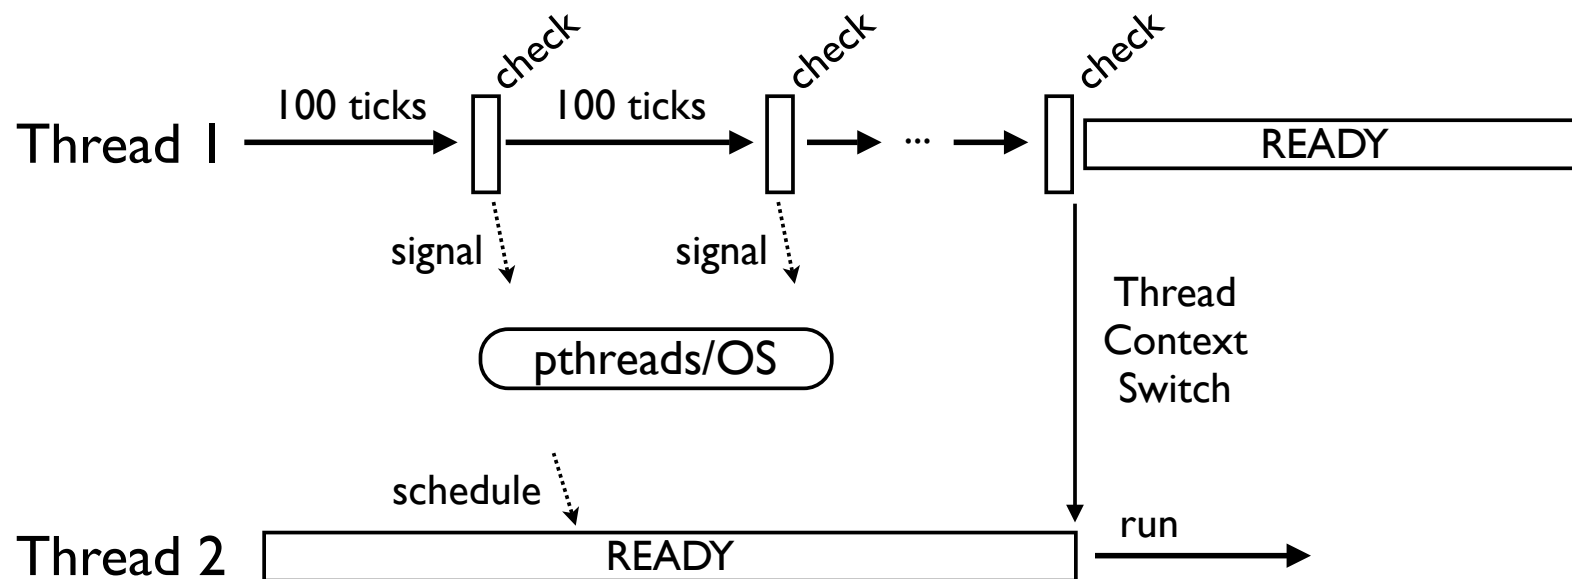
- Trace files were large (>20MB for 1s of running)

# Logging Results

- The logs were quite revealing
- Interesting behavior on one CPU
- Diabolical behavior on multiple CPUs
- Will briefly summarize findings followed by an interactive visualization that shows details

# Single CPU Threading

- Threads alternate execution, but switch far less frequently than you might imagine




- Hundreds to thousands of checks might occur before a thread context switch (this is good)



# Multicore GIL War

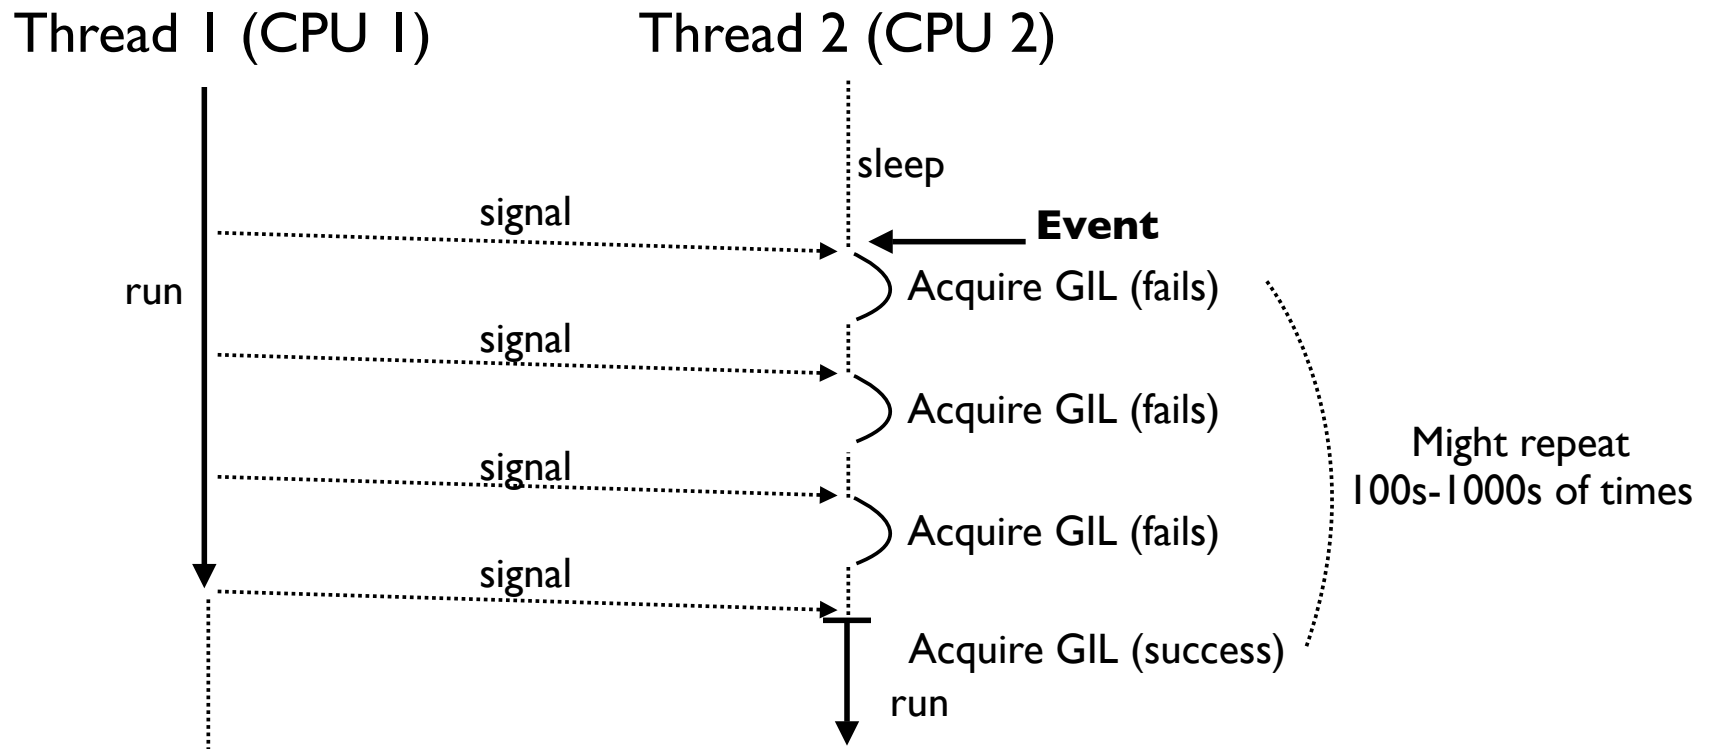
- With multiple cores, runnable threads get scheduled simultaneously (on different cores) and battle over the GIL



- Thread 2 is repeatedly signaled, but when it wakes up, the GIL is already gone (reacquired)


# Multicore Event Handling

- CPU-bound threads make GIL acquisition difficult for threads that want to handle events



# Behavior of I/O Handling

- I/O ops often do not block



- Due to buffering, the OS is able to fulfill I/O requests immediately and keep a thread running
- However, the GIL is always released
- Results in GIL thrashing under heavy load

# GIL Visualization (Demo)

- Let's look at all of these effects

<http://www.dabeaz.com/GIL>

- Some facts about the plots:
  - Generated from ~2GB of log data
  - Rendered into ~2 million PNG image tiles
  - Created using custom scripts/tools
  - I used the multiprocessing module

# Part 4

## A Better GIL?

# The New GIL

- Python 3.2 has a new GIL implementation (only available by svn checkout)
- The work of Antoine Pitrou (applause)
- It aims to solve all that GIL thrashing
- It is the first major change to the GIL since the inception of Python threads in 1992
- Let's go take a look

# New Thread Switching

- Instead of ticks, there is now a global variable


```
/* Python/ceval.c */  
...
```

```
static volatile int gil_drop_request = 0;
```

- A thread runs until the value gets set to 1
- At which point, the thread must drop the GIL
- Big question: How does that happen?

# New GIL Illustrated

- Suppose that there is just one thread


Thread 1 


- It just runs and runs and runs ...
  - Never releases the GIL
  - Never sends any signals
  - Life is great!



# New GIL Illustrated

- Suppose, a second thread appears


Thread 1 


Thread 2 

- It is suspended because it doesn't have the GIL
- Somehow, it has to get it from Thread 1

# New GIL Illustrated

- Waiting thread does a timed `cv_wait` on GIL

Thread 1 

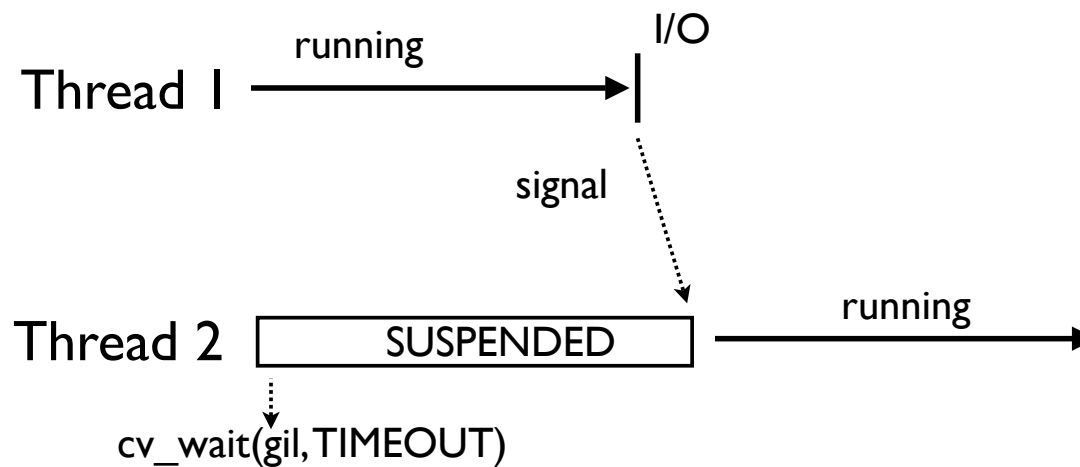
Thread 2   
↓  
`cv_wait(gil, TIMEOUT)`

By default `TIMEOUT`  
is 5 milliseconds, but it  
can be changed

- The idea : Thread 2 waits to see if the GIL gets released voluntarily by Thread 1 (e.g., if there is I/O or it goes to sleep for some reason)

# New GIL Illustrated

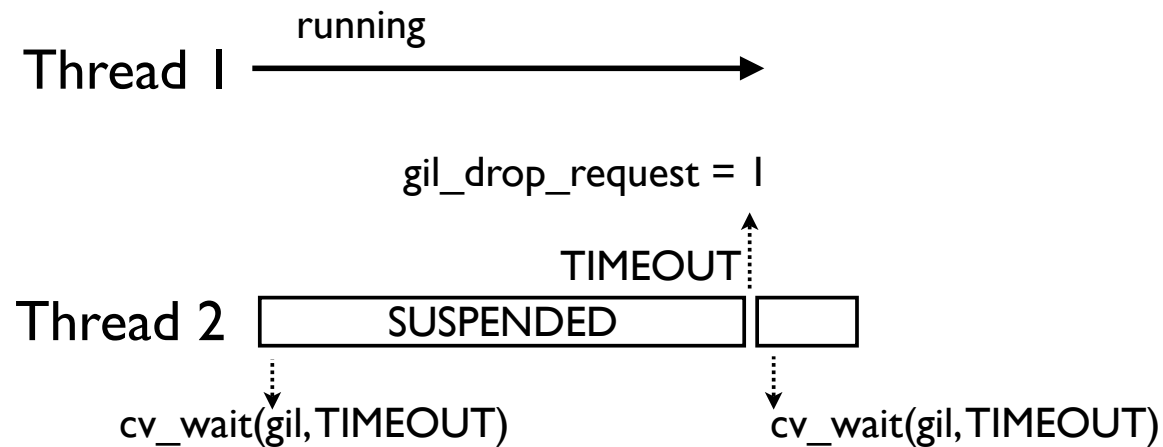
- Voluntary GIL release



- This is the easy case. Second thread is signaled and it grabs the GIL.

# New GIL Illustrated

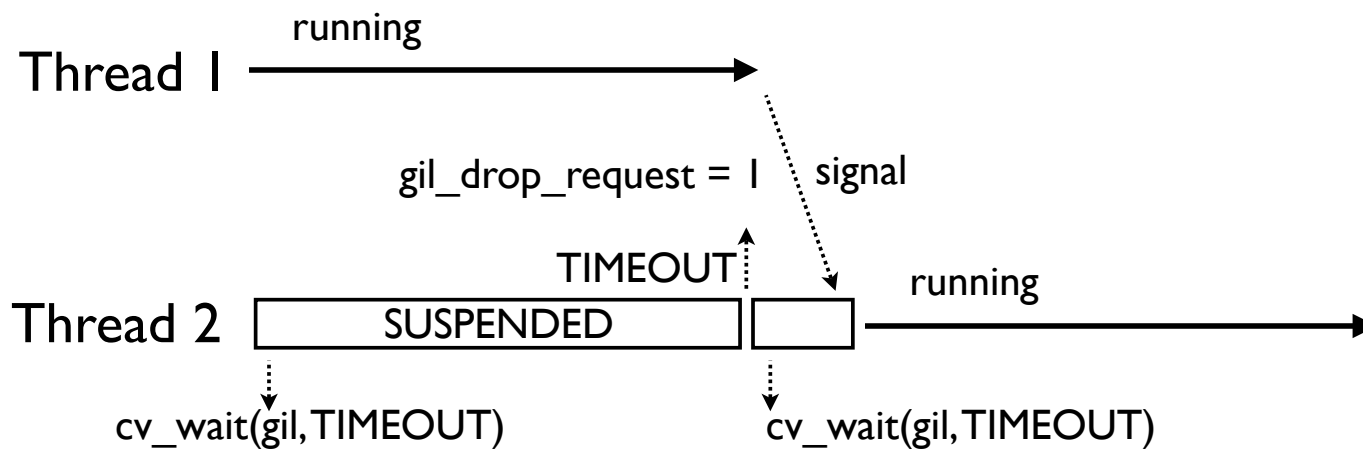
- If timeout, set `gil_drop_request`



- Thread 2 then repeats its wait on the GIL

# New GIL Illustrated

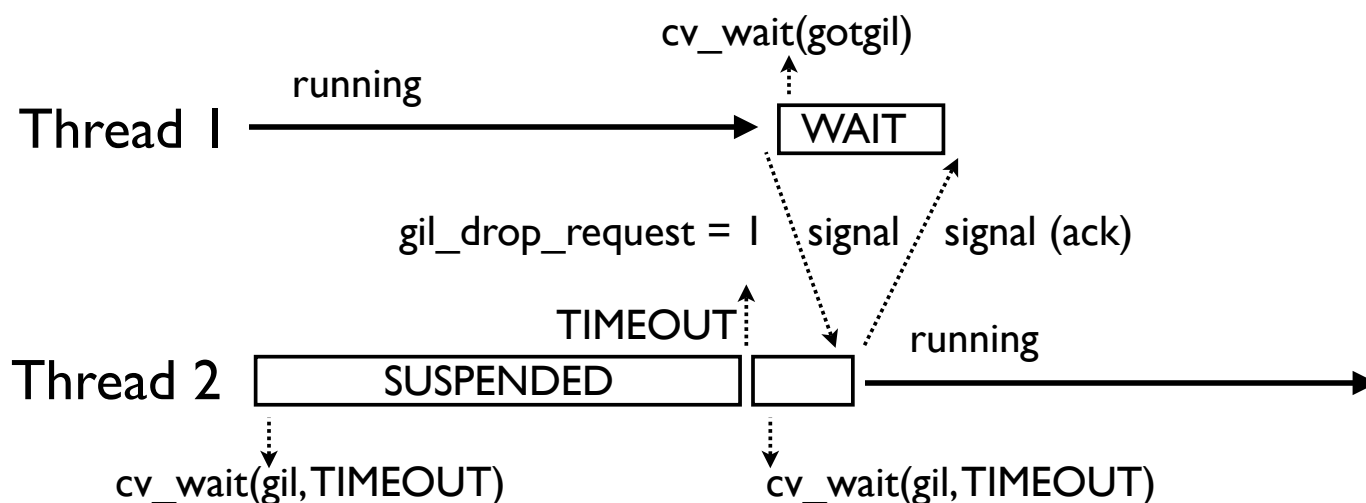
- Thread 1 suspends after current instruction



- Signal is sent to indicate release of GIL

# New GIL Illustrated

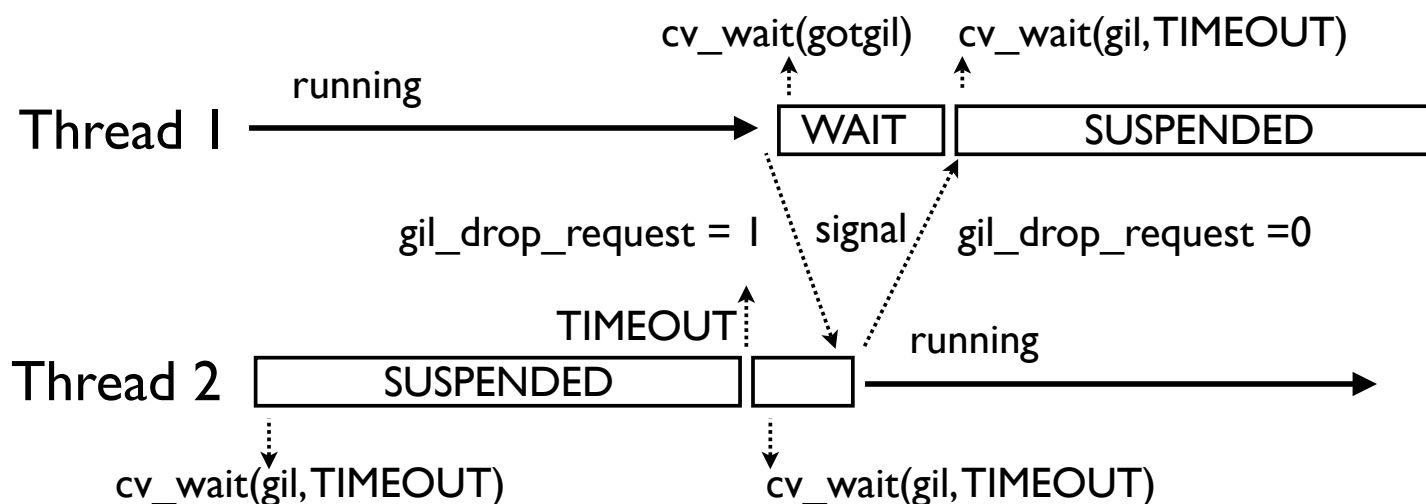
- On a forced release, a thread waits for an ack



- Ack ensures that the other thread successfully got the GIL and is now running
- This eliminates the "GIL Battle"

# New GIL Illustrated

- The process now repeats itself for Thread 1



- So, the timeout sequence happens over and over again as CPU-bound threads execute

# Does it Work?

- Yes, apparently (4-core MacPro, OS-X 10.6.2)

Sequential : 11.53s

Threaded (2 threads) : 11.93s

Threaded (4 threads) : 12.32s

- Keep in mind, Python is still limited by the GIL in all of the usual ways (threads still provide no performance boost)
- But, otherwise, it looks promising!



# Part 5

## Die GIL Die!!!

# Alas, It Doesn't Work

- The New GIL impacts I/O performance
- Here is a fragment of network code

## Thread 1

```
def spin():  
    while True:  
        # some work  
        pass
```

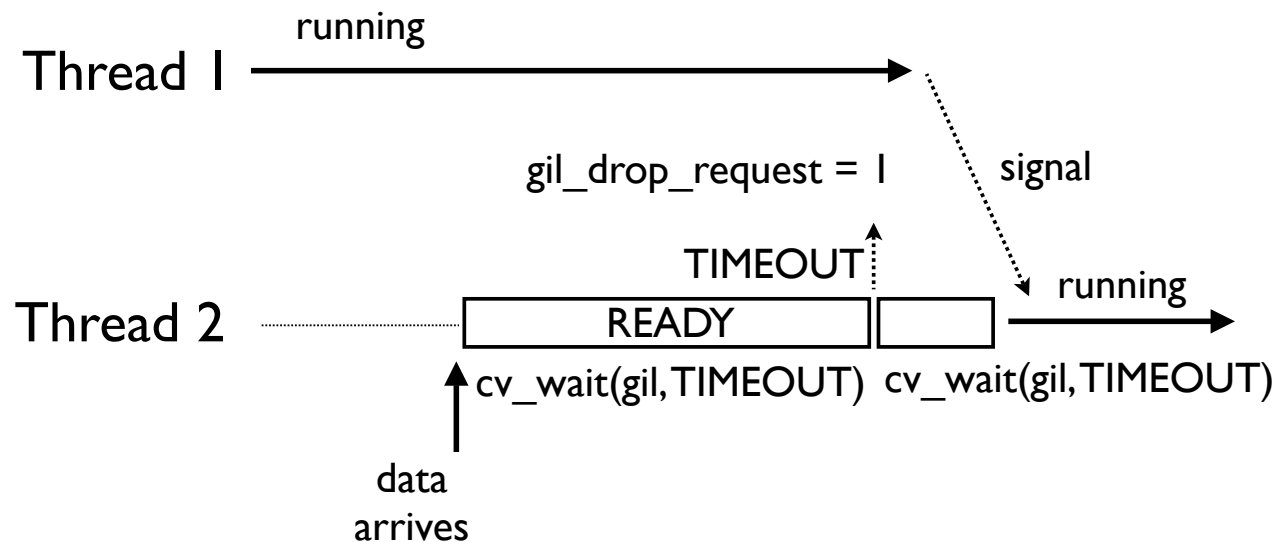
## Thread 2

```
def echo_server(s):  
    while True:  
        data = s.recv(8192)  
        if not data:  
            break  
        s.sendall(data)
```

- One thread is working (CPU-bound)
- One thread receives and echos data on a socket

# Response Time

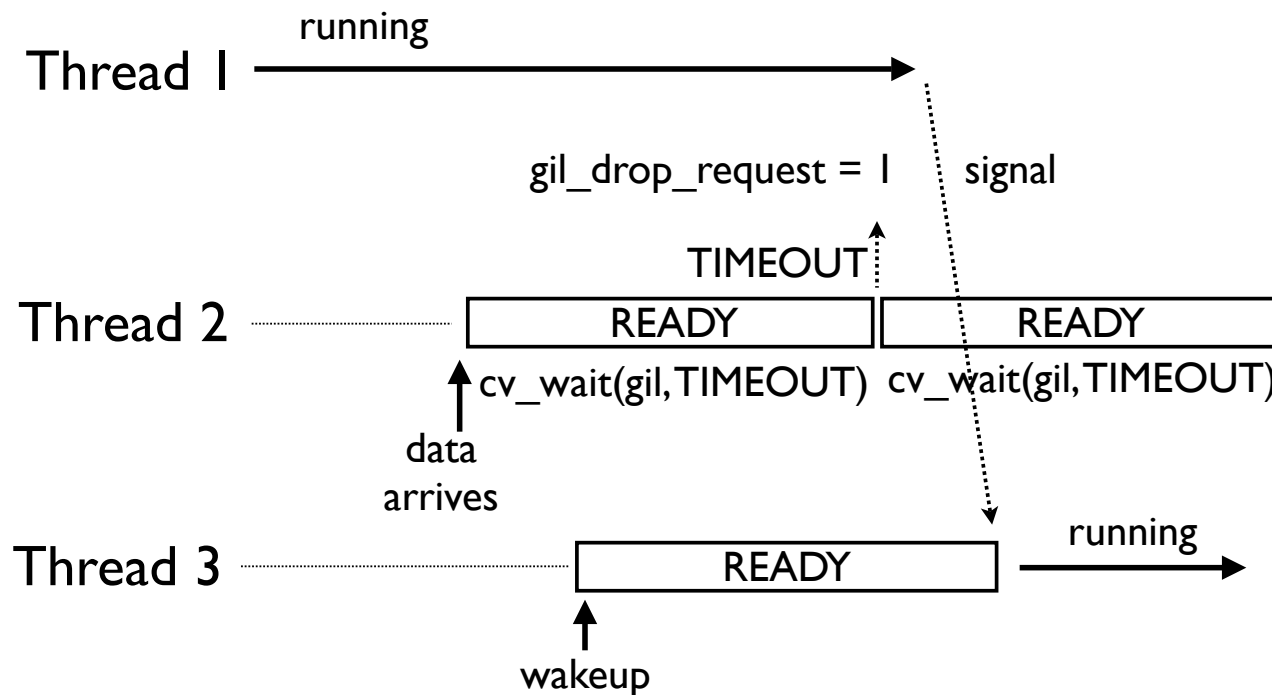
- New GIL increases response time



- To handle I/O, a thread must go through the entire timeout sequence to get control
- Ignores the high priority of I/O or events

# Unfair Wakeup/Starvation

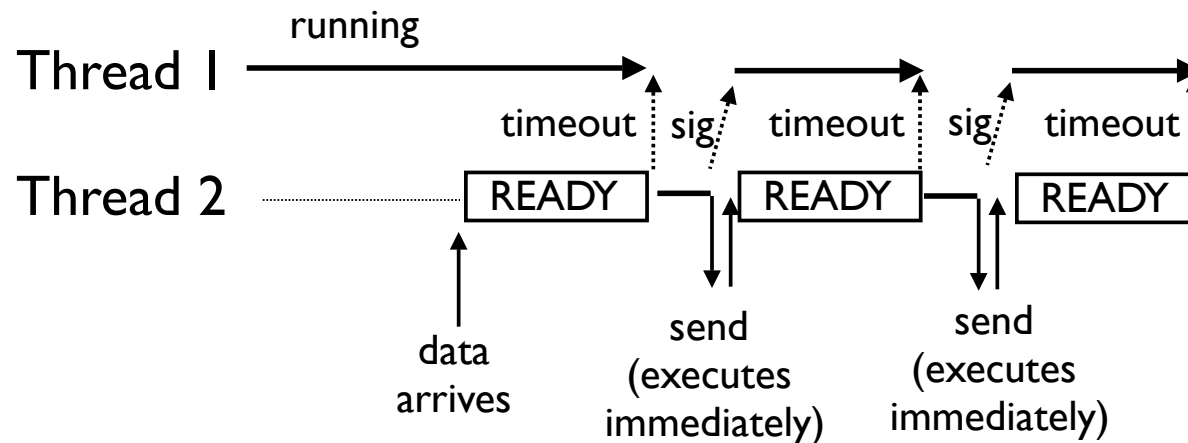
- Most deserving thread may not get the GIL



- Caused by internal condition variable queuing
- Further increases the response time

# Convoy Effect

- I/O operations that don't block cause stalls



- Since I/O operations always release the GIL, CPU-bound threads will always try to restart
- On I/O completion (almost immediately), the GIL is gone so the timeout has to repeat

# An Experiment

- Send 10MB of data to an echo server thread that's competing with a CPU-bound thread

Python 2.6.4 (2 CPU) : 0.57s (10 sample average)

Python 3.2 (2 CPU) : 12.4s (20x slower)

- What if echo competes with 2 CPU threads?

Python 2.6.4 (2 CPU) : 0.25s (Better performance?)

Python 3.2 (2 CPU) : 46.9s (4x slower than before)

Python 3.2 (1 CPU) : 0.14s (330x faster than 2 cores?)

- Arg! Enough already!

# Part 6

Score : Multicore 2, GIL 0

# Fixing the GIL

- Can the GIL's erratic behavior be fixed?
- My opinion :Yes, maybe.
- The new GIL is already 90% there
- It just needs a few extra bits



# The Missing Bits

- Priorities: There must be some way to separate CPU-bound (low priority) and I/O bound (high priority) threads
- Preemption: High priority threads must be able to immediately preempt low-priority threads

# A Possible Solution

- Operating systems use timeouts to automatically adjust task priorities (multilevel feedback queuing)
  - If a thread is preempted by a timeout, it is penalized with lowered priority (bad thread)
  - If a thread suspends early, it is rewarded with raised priority (good thread)
  - High priority threads always preempt low priority threads
- Maybe it could be applied to the new GIL?

# Remove the GIL?

- This entire talk has been about the problem of implementing one tiny little itty bitty lock
- Fixing Python to remove the GIL entirely is an exponentially more difficult project
- If there is one thing to take away, there are practical reasons why the GIL remains

# Final Thoughts

- Don't use this talk to justify not using threads
- Threads are a very useful programming tool for many kinds of concurrency problems
- Threads can also offer excellent performance even with the GIL (you need to study it)
- However, you should know about the tricky corner cases

# Final Thoughts

- Improving the GIL is something that all Python programmers should care about
- Multicore is not going away
- You might not use threads yourself, but they are used for a variety of low-level purposes in frameworks and libraries you might be using
- More predictable thread behavior is good

# Thread Terminated

- That's it!
- Thanks for listening!
- Questions